
Virtual Memory

Philipp Koehn

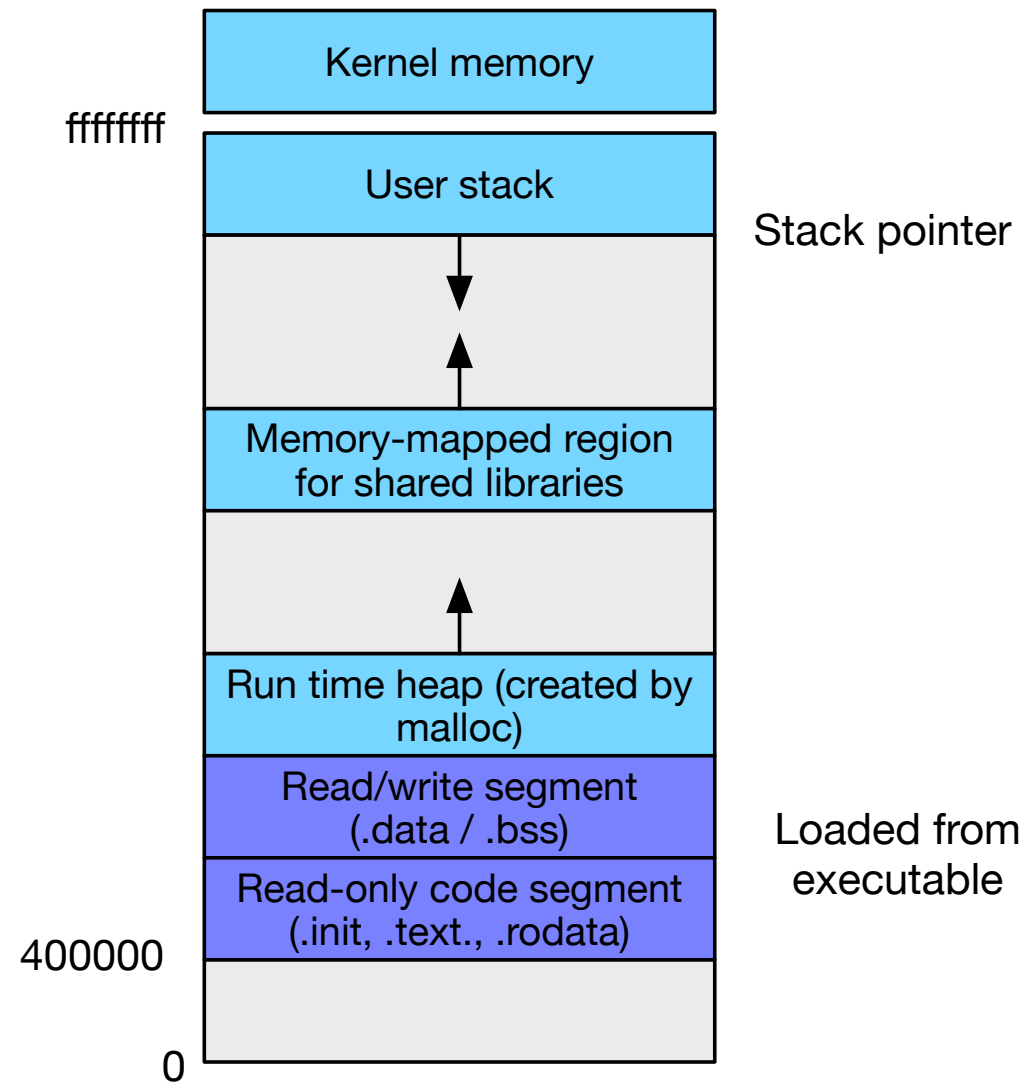
25 April 2018



Recall: Process Address Space



1



Virtual Memory



2

- Abstraction of physical memory
- Purpose
 - appearance of more available memory than physically exists (DRAM)
 - handles disk caching / loading
 - insulates memory of each process

Virtual Memory



2

- Abstraction of physical memory
- Purpose
 - appearance of more available memory than physically exists (DRAM)
 - handles disk caching / loading
 - insulates memory of each process
- Page table: maps from virtual address to physical addresses

Virtual Memory



- Abstraction of physical memory
- Purpose
 - appearance of more available memory than physically exists (DRAM)
 - handles disk caching / loading
 - insulates memory of each process
- Page table: maps from virtual address to physical addresses
- Memory management unit (MMU):
hardware implementation of address translation


Warning



- This is going to get very complex
- Closely tied with multi-tasking (multiple processes)
- Partly managed by hardware,
partly managed by software

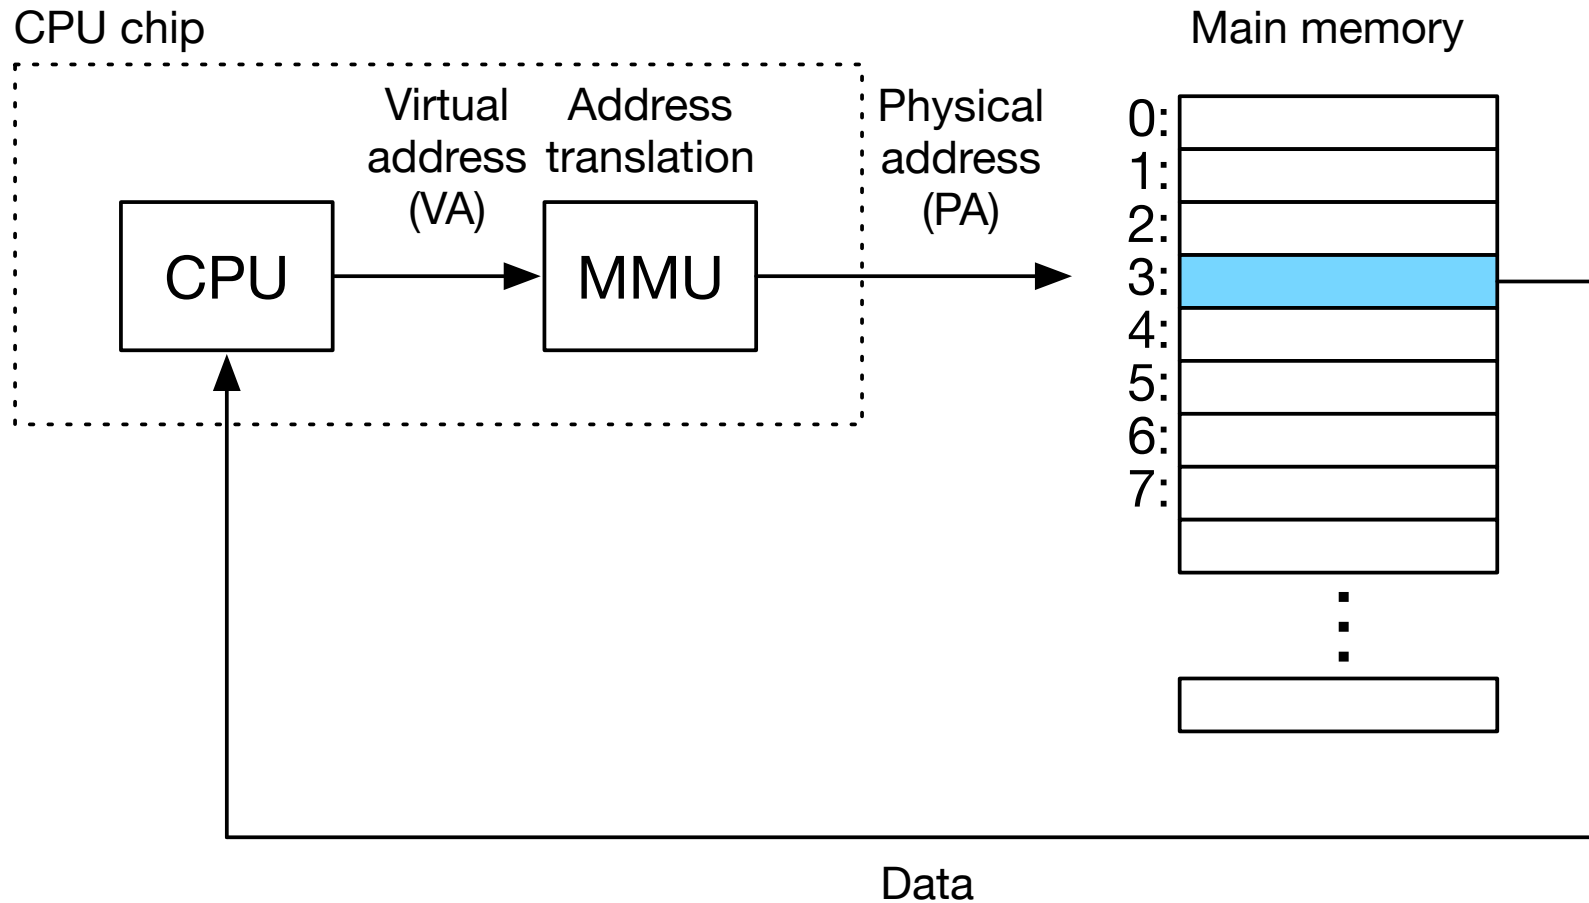
virtual addressing

Physical Addressing



- So far, assumed CPU addresses physical memory

Virtual Addressing



- Memory management unit (MMU): maps virtual to physical addresses

Address Space



- Virtual memory size: $N = 2^n$ bytes, e.g., 256TB
- Physical memory size: $M = 2^m$ bytes, e.g., 16GB
- Page (block of memory): $P = 2^p$ bytes, e.g., 4KB
- A virtual address can be encoded in n bits

caching

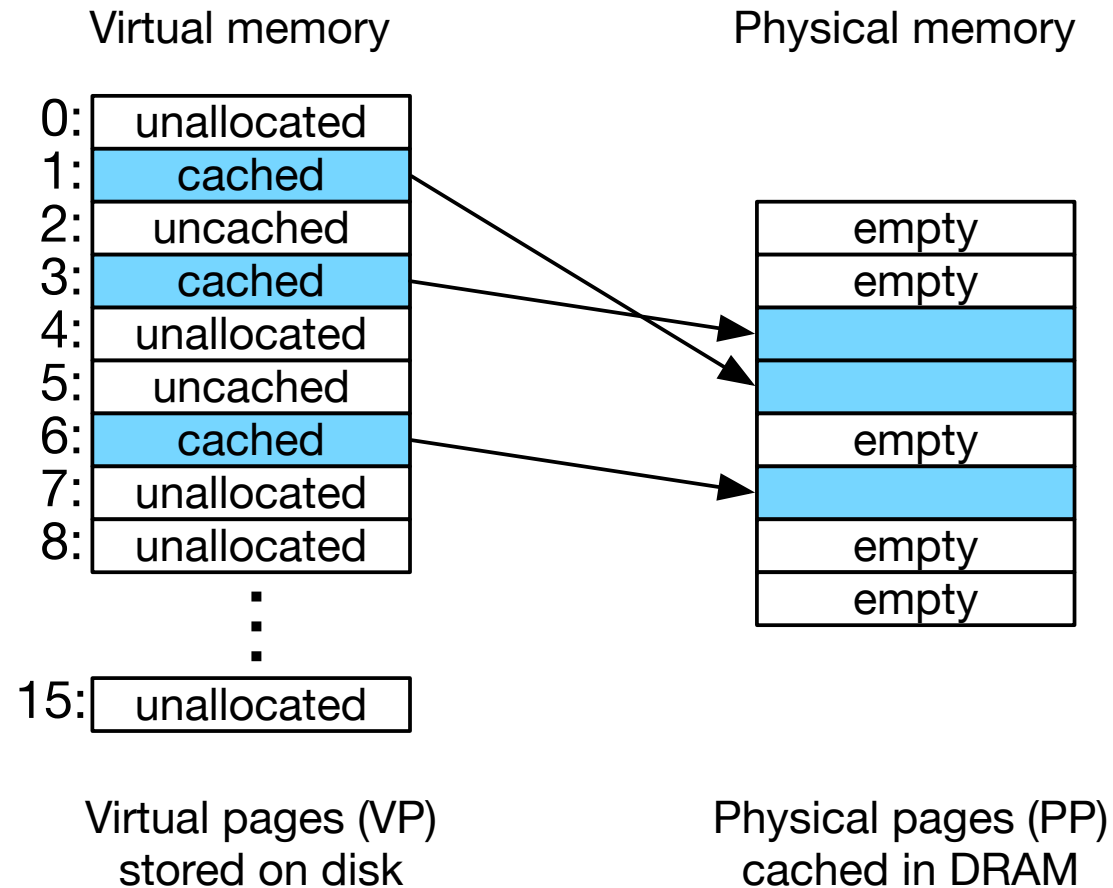
Caching... Again?



9

- Yes, we already discussed caching
but: for on-chip cache of DRAM memory
- Now
 - caching between RAM and disk
 - driven by a large virtual memory address space
 - to avoid unnecessary and duplicate loading
- Jargon
 - previously "block", now "page"
 - now: "swapping" or "paging"

Mapping



State of Virtual Memory Page

11



- Cached
 - allocated page
 - stored in physical memory

State of Virtual Memory Page

- Cached
 - allocated page
 - stored in physical memory
- Uncached
 - allocated page
 - not in physical memory

State of Virtual Memory Page

- Cached
 - allocated page
 - stored in physical memory
- Uncached
 - allocated page
 - not in physical memory
- Unallocated
 - not used by virtual memory system so far

Page Table

12



- Array of page table entries (PTE)

Page Table

12



- Array of page table entries (PTE)
- Valid bit
 - set if PTE currently maps to physical address (cached)
 - not set otherwise (uncached or unallocated)

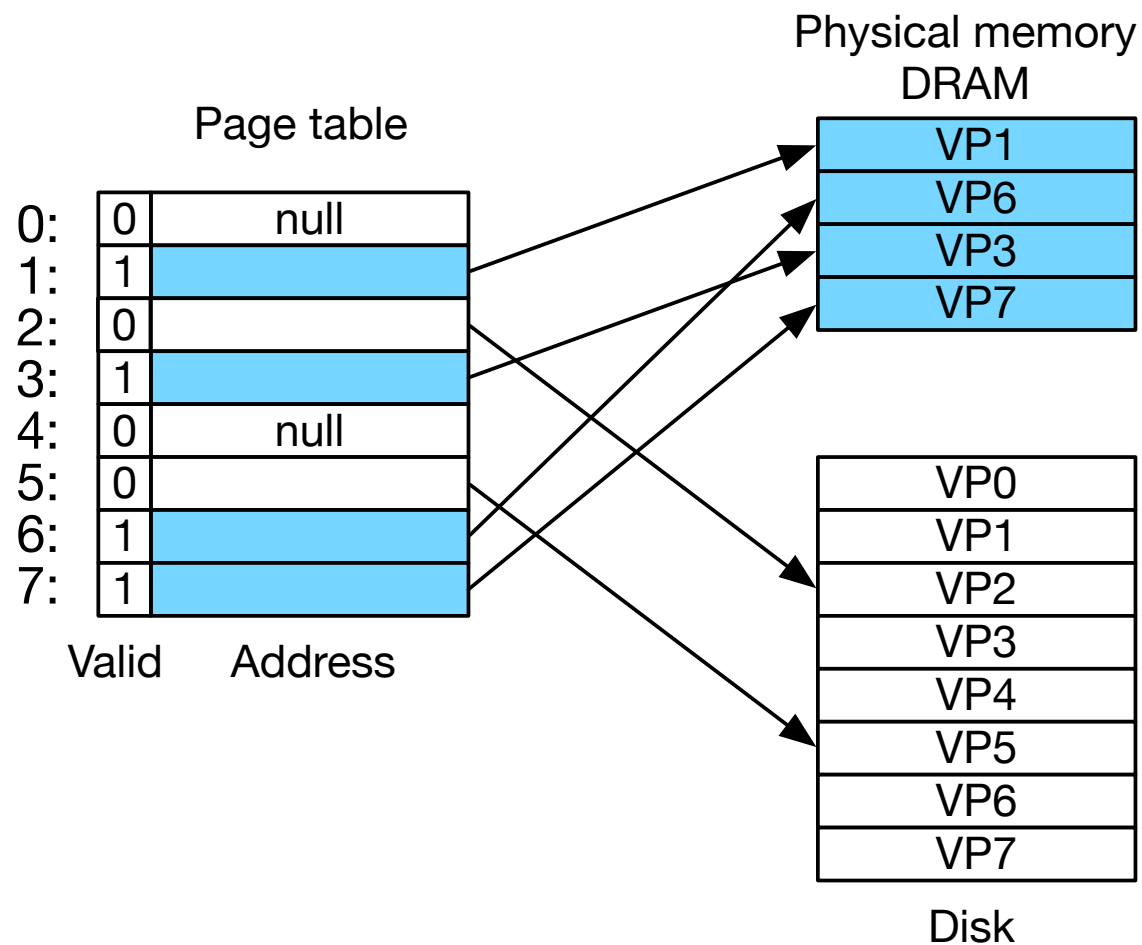
Page Table

12

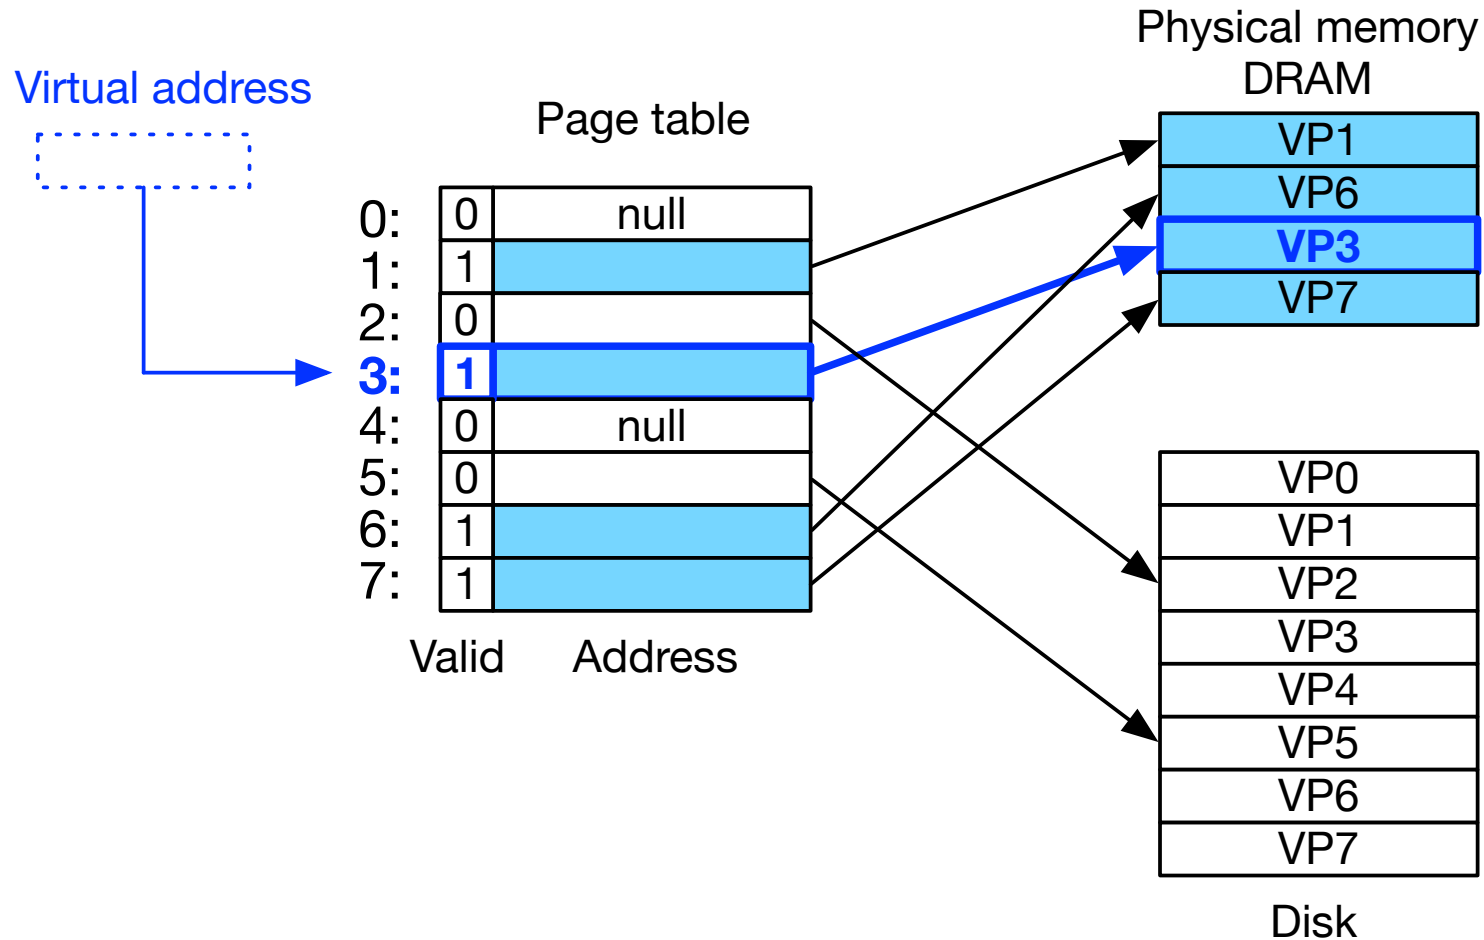


- Array of page table entries (PTE)
- Valid bit
 - set if PTE currently maps to physical address (cached)
 - not set otherwise (uncached or unallocated)
- Mapped address
 - if cached: physical address in DRAM
 - if not cached: physical address on disk

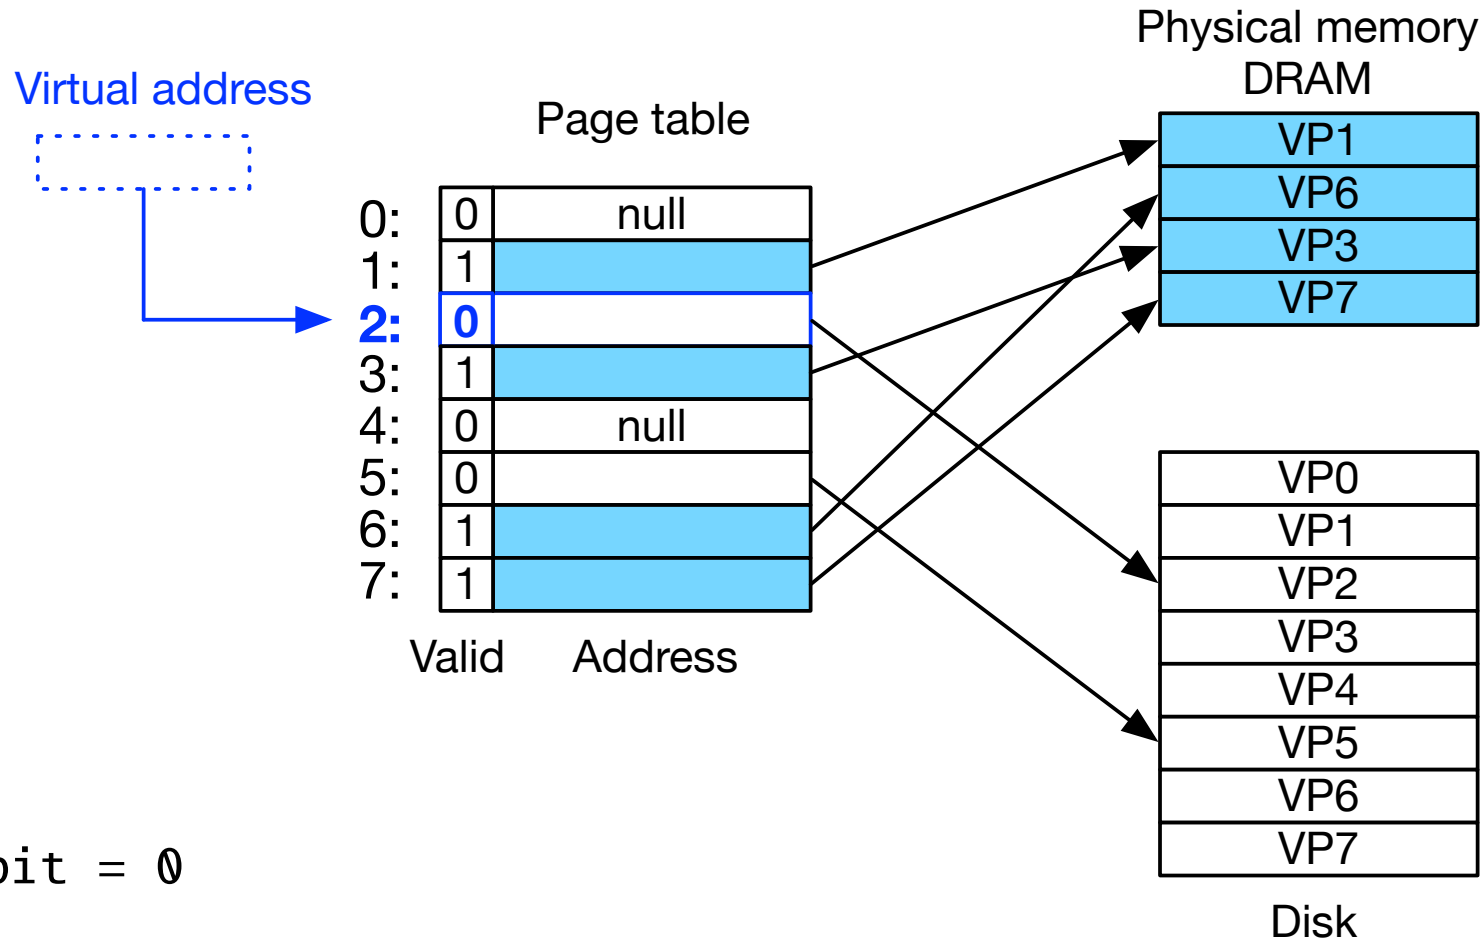
Page Table



Page Hit

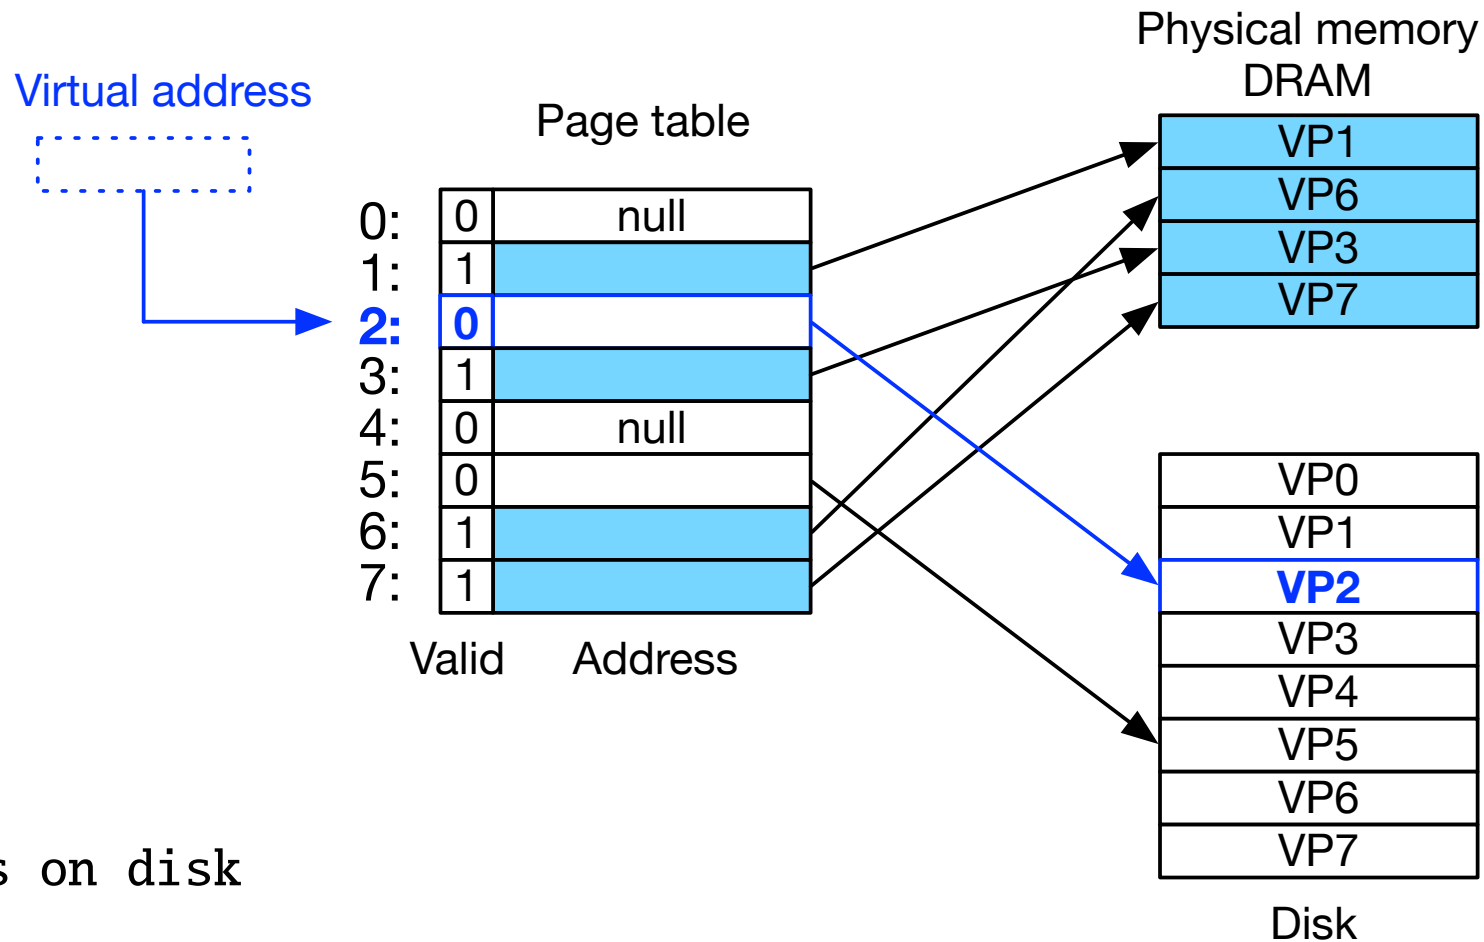


Page Fault



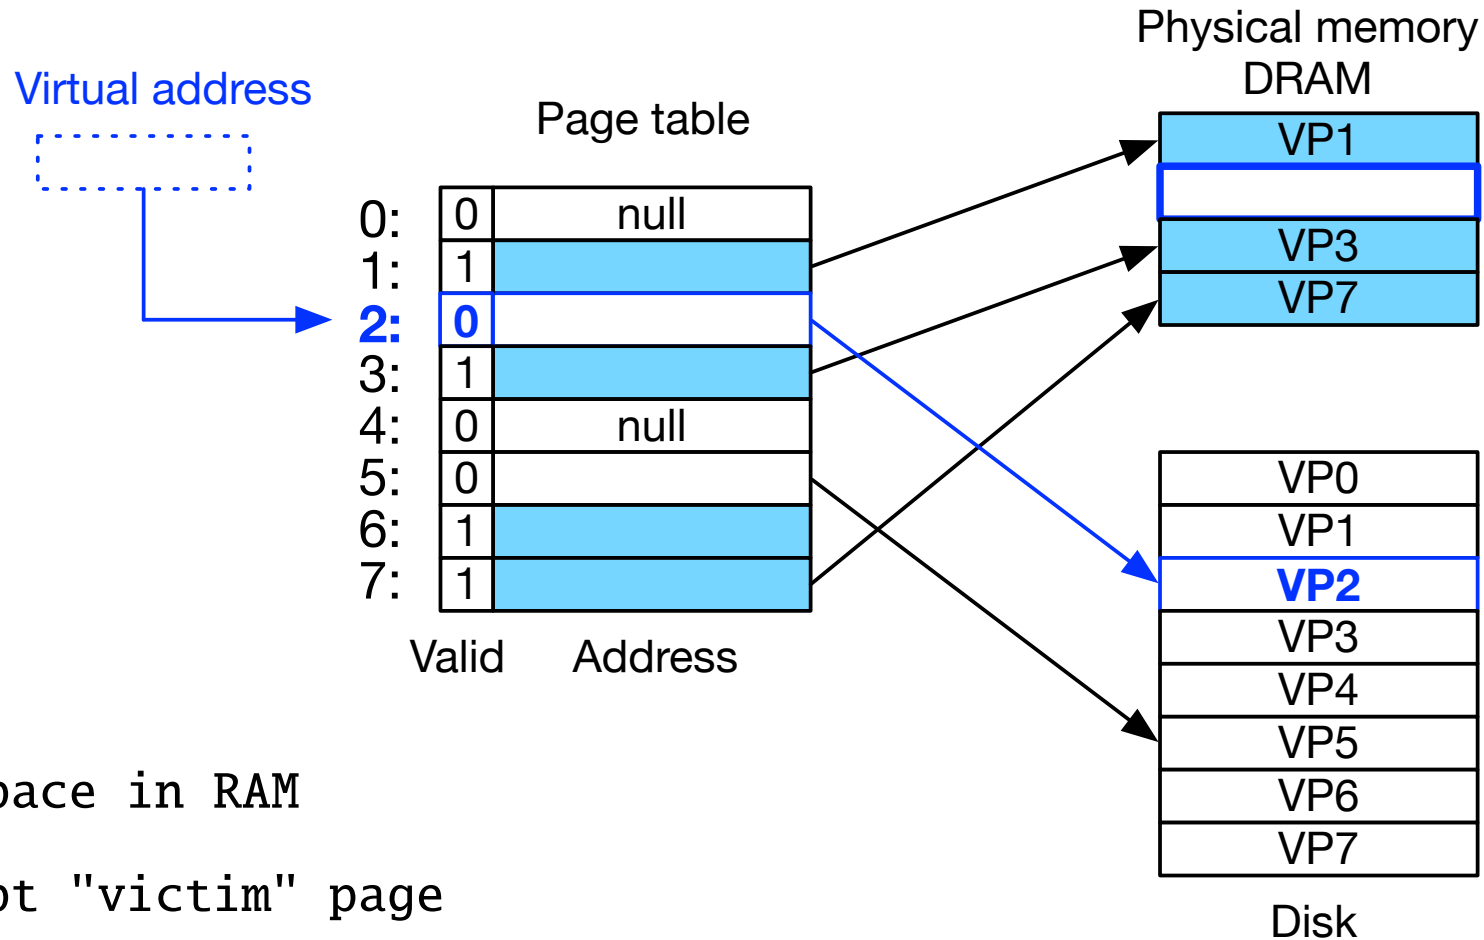
- Valid bit = 0
- Page not in RAM

Page Fault



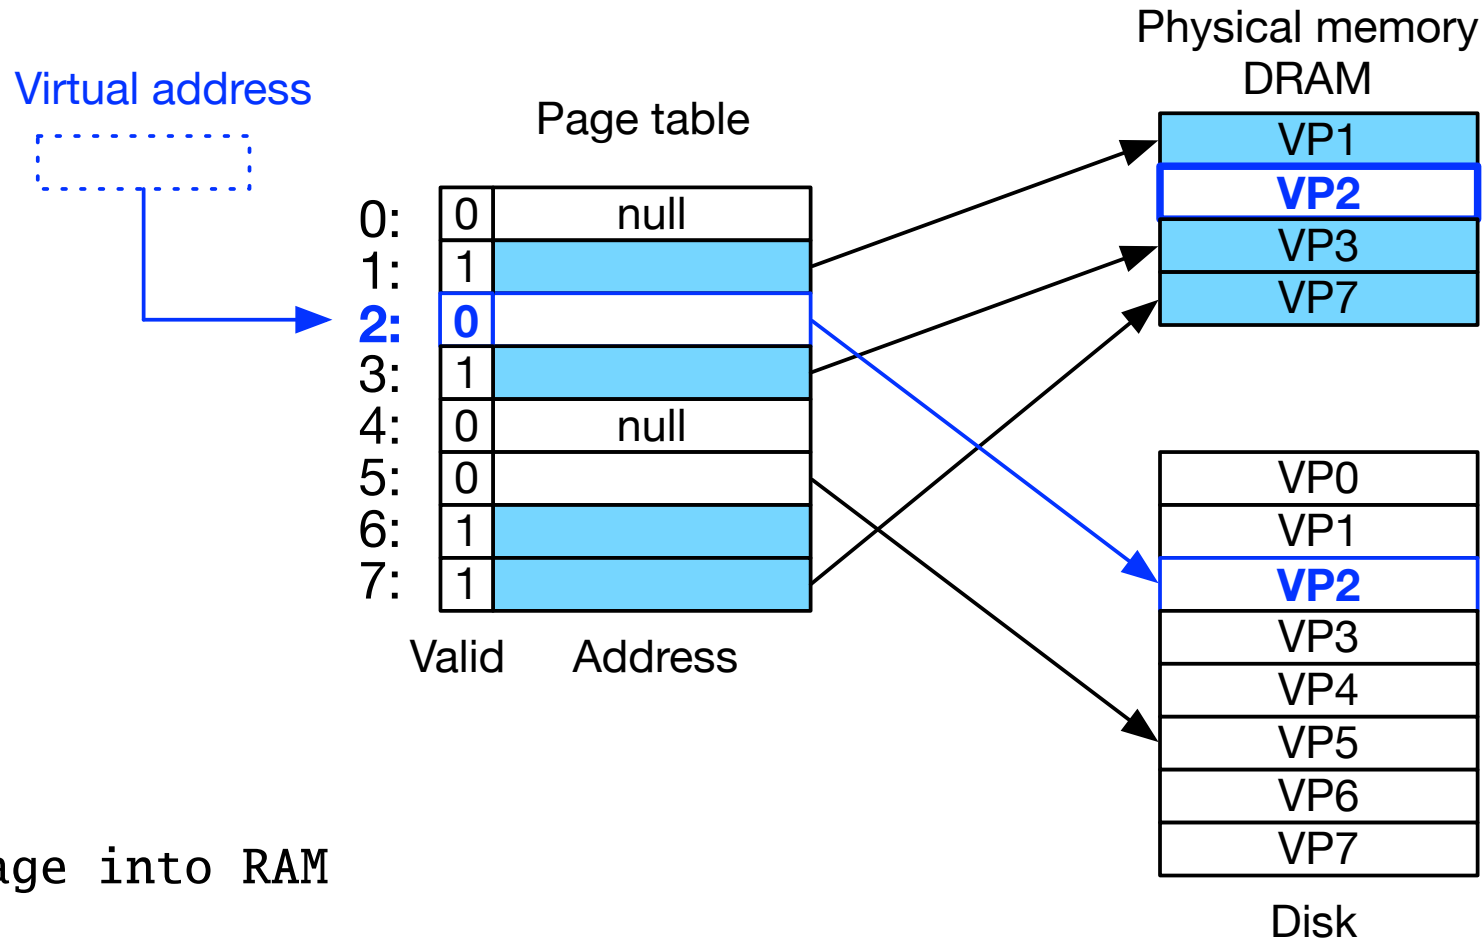
- Page is on disk

Page Fault



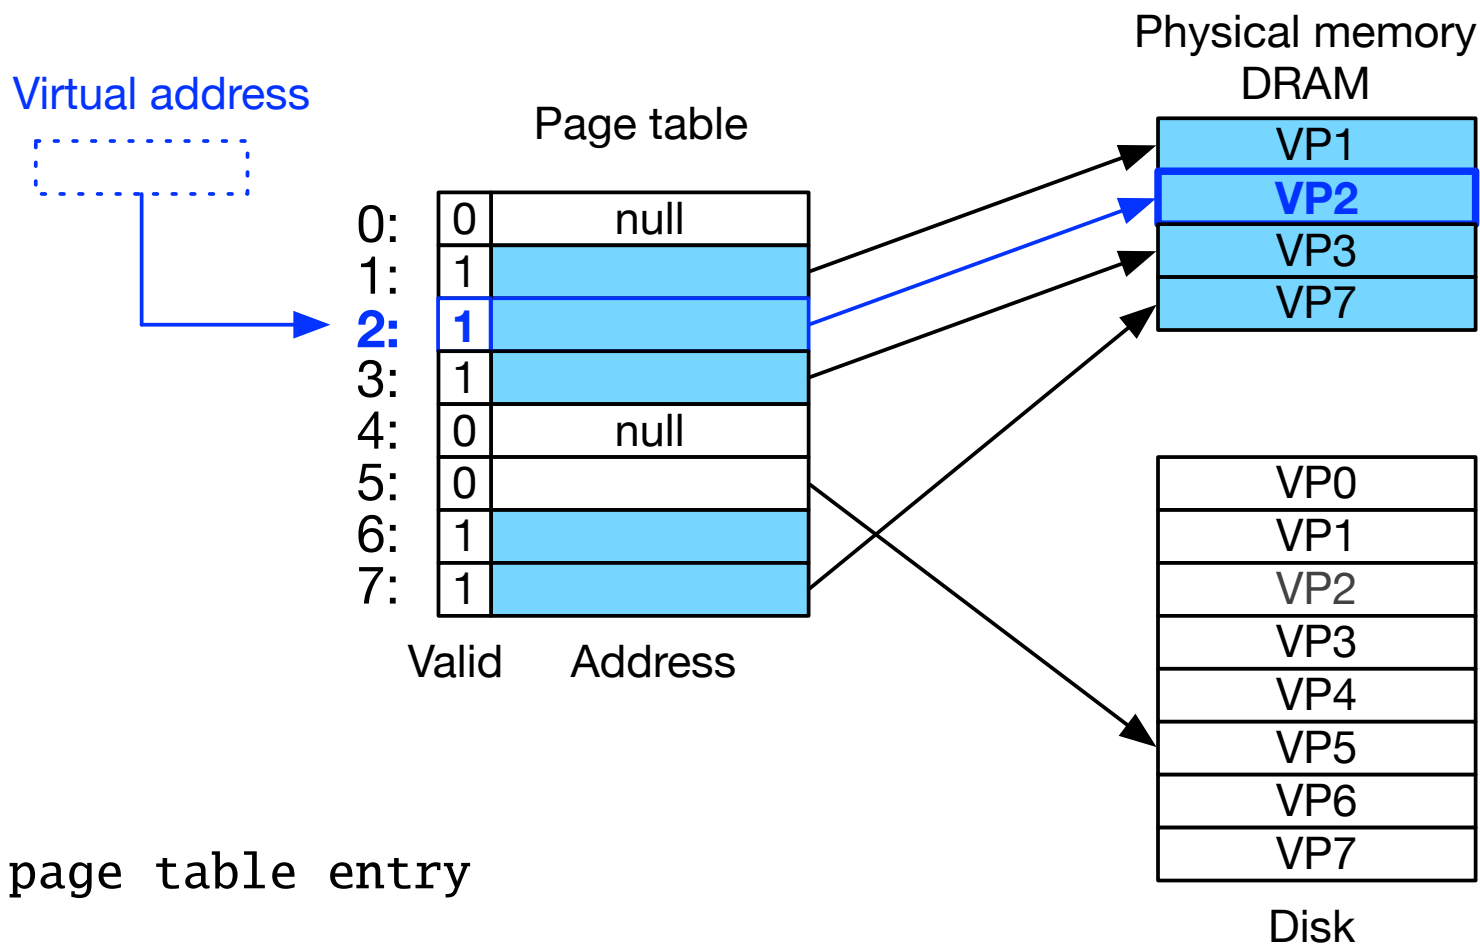
- Make space in RAM
- Pre-empt "victim" page
- Typically out-dated cached page

Page Fault



- Load page into RAM

Page Fault

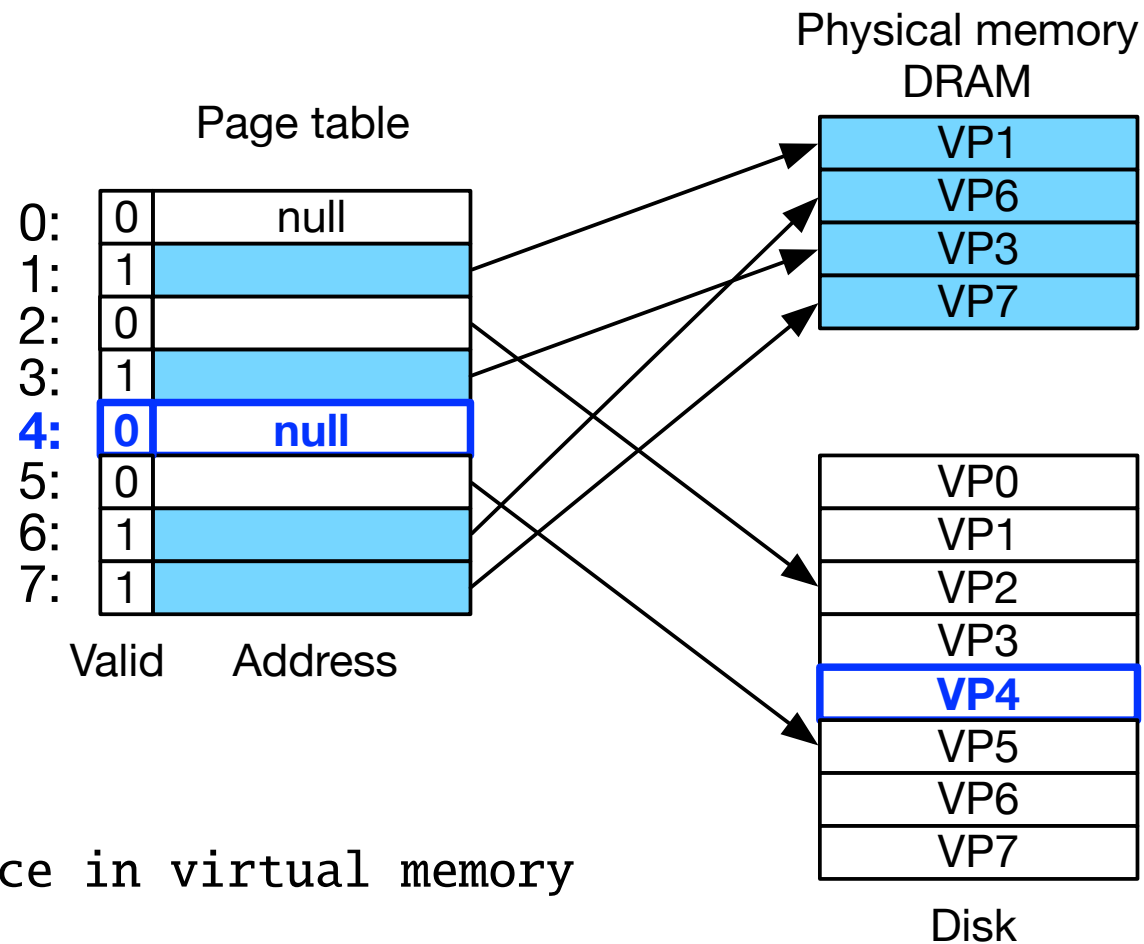


- Update page table entry

Allocating Pages

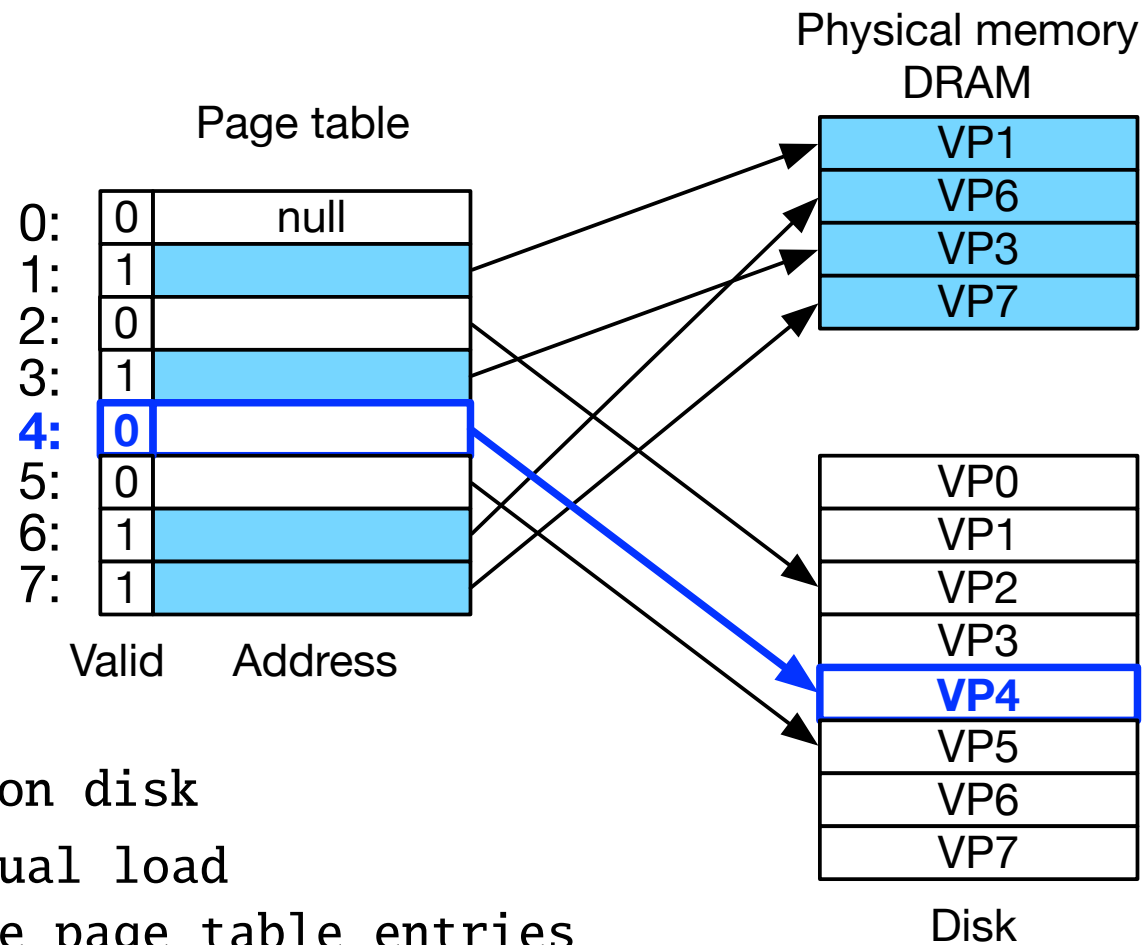
- What happens when we load a program?
- We need to load its executable into memory
- Similar: create data objects when program is running ("allocating" memory)

Allocating Page



- Identify space in virtual memory

Allocating Page



- Map to data on disk
 - do not actual load
 - just create page table entries
 - let virtual memory system handle loading

⇒ On-demand loading

Process Memory

- Nothing loaded at startup

Process Memory

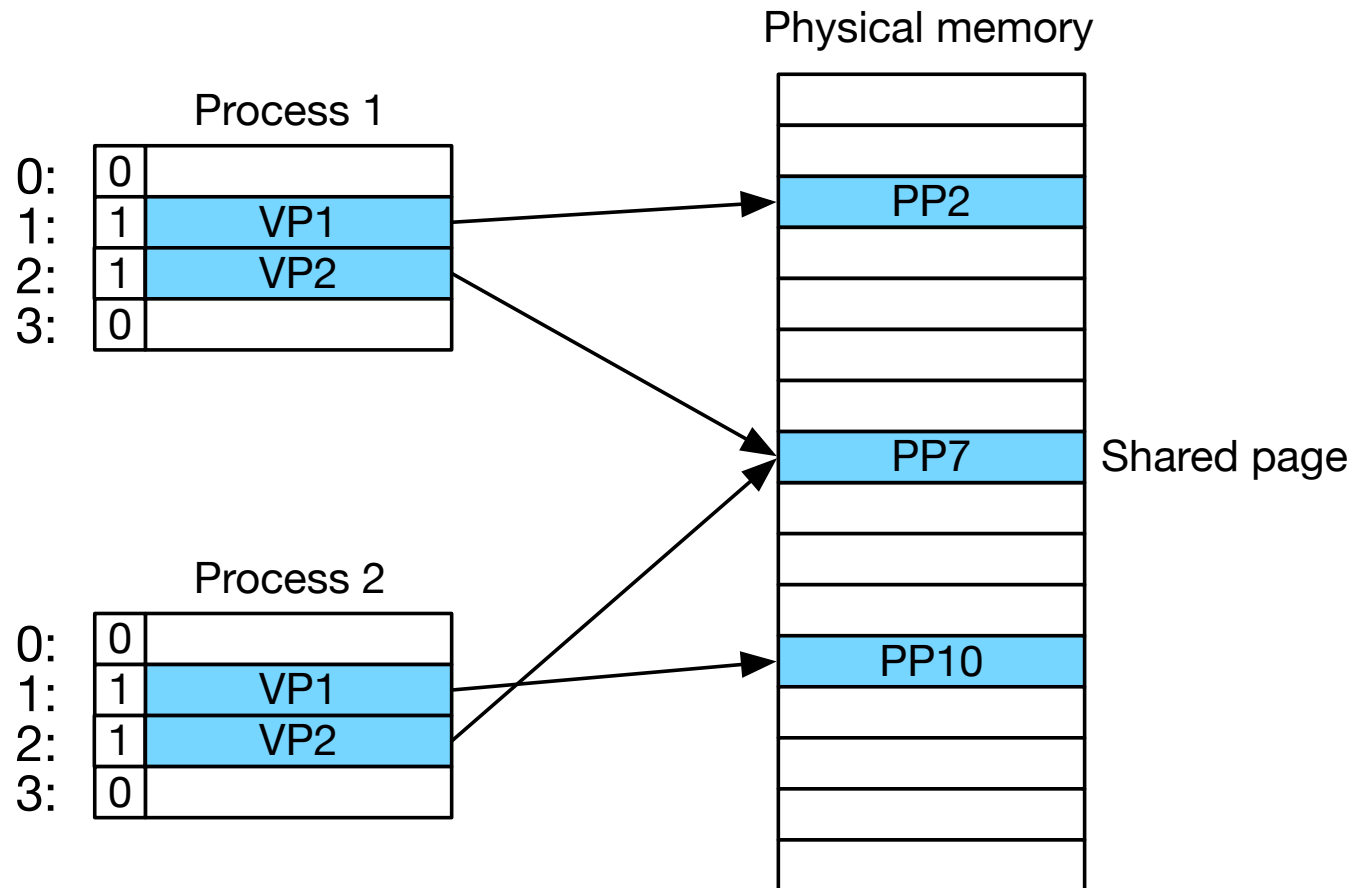
- Nothing loaded at startup
- Working set (or resident set)
 - pages of a process that are currently in DRAM
 - loaded by virtual memory system on demand

- Nothing loaded at startup
- Working set (or resident set)
 - pages of a process that are currently in DRAM
 - loaded by virtual memory system on demand
- Thrashing
 - memory actively required by all processes larger than physically available
 - frequent swapping of memory to/from disk
 - very bad: slows down machine dramatically

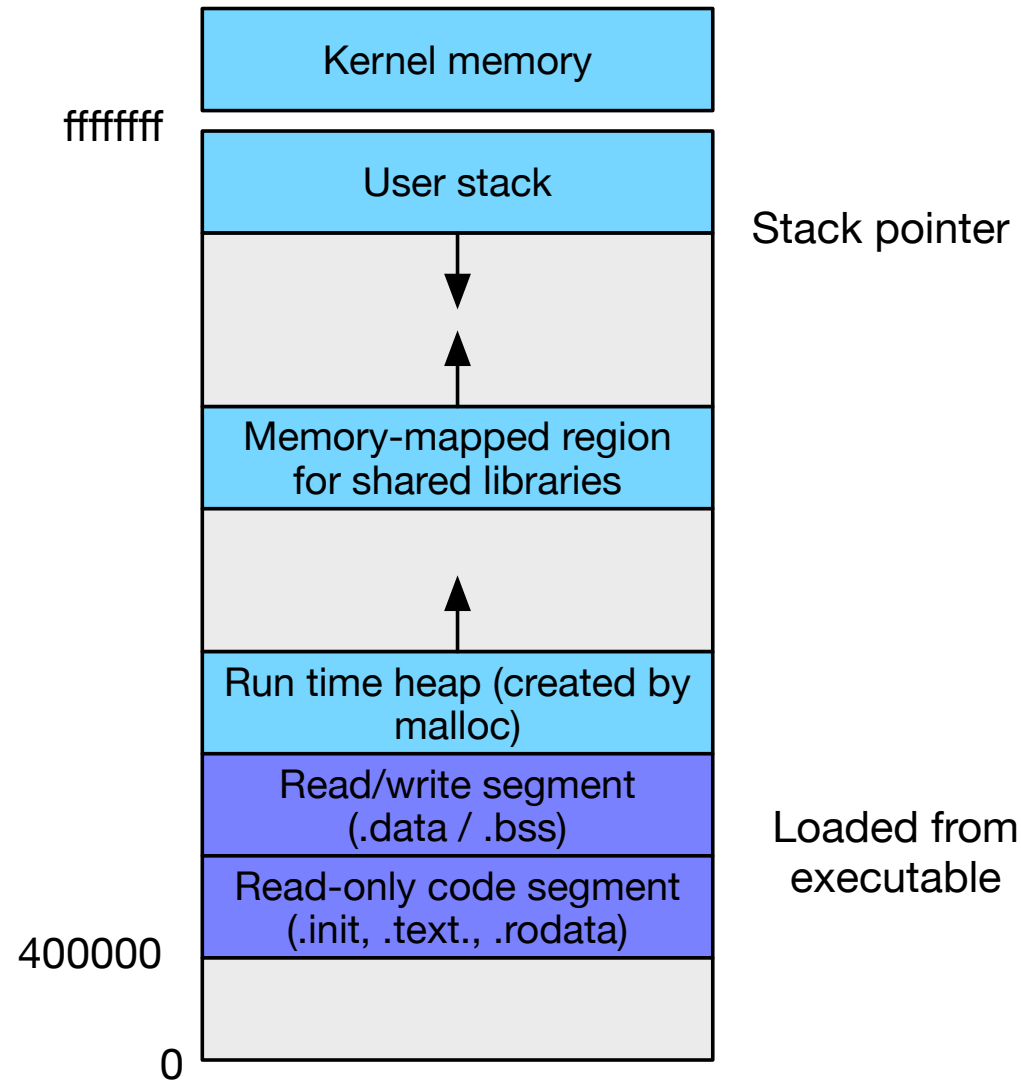


memory management


One Page Table per Process



Process Address Space



Simplified Linking



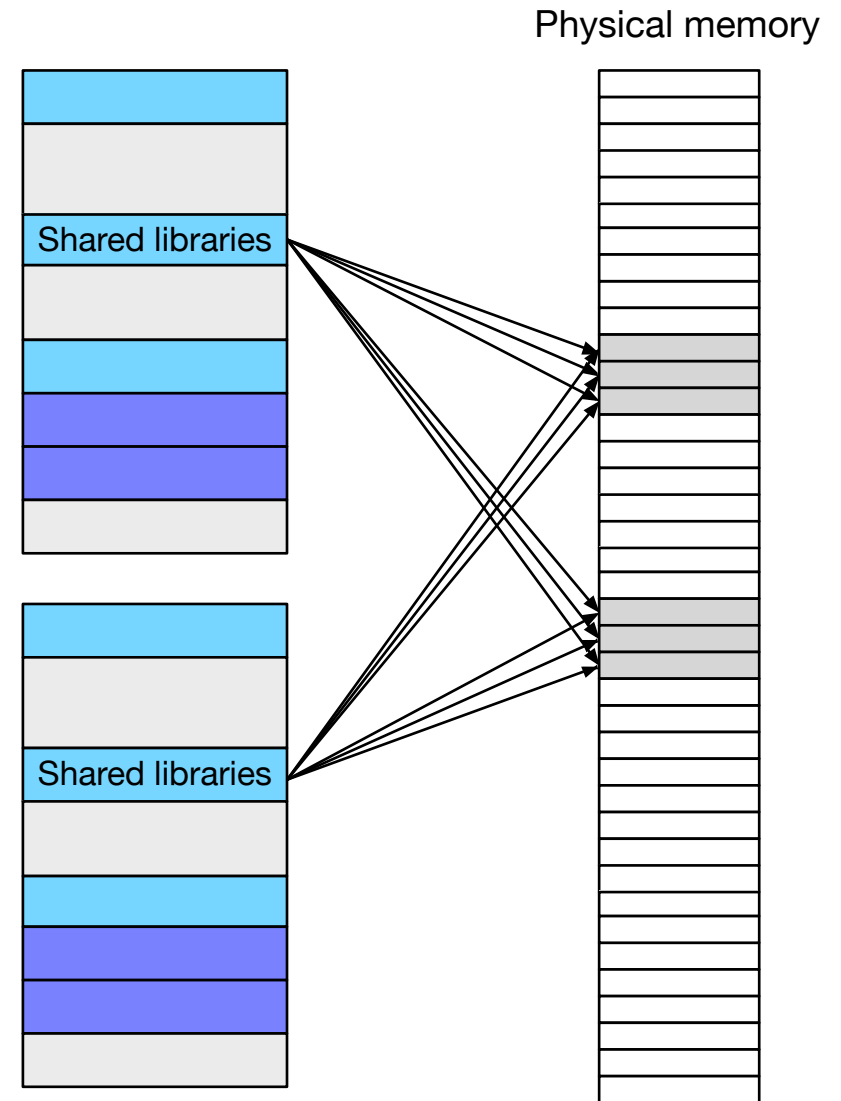
- Each process has its code in address `0x400000`
- Easy linking: Linker can establish fixed addresses

Simplified Loading

- When loading process into memory...
- Enter `.data` and `.text` section into page table
- Mark them as invalid (= not actually in RAM)
- Called memory mapping (more on that later)

Simplified Sharing

- Shared libraries used by several processes
e.g., `stdio` providing `printf`, `scanf`, `open`, `close`, ...
- Not copied multiple times into RAM



Simplified Memory Allocation

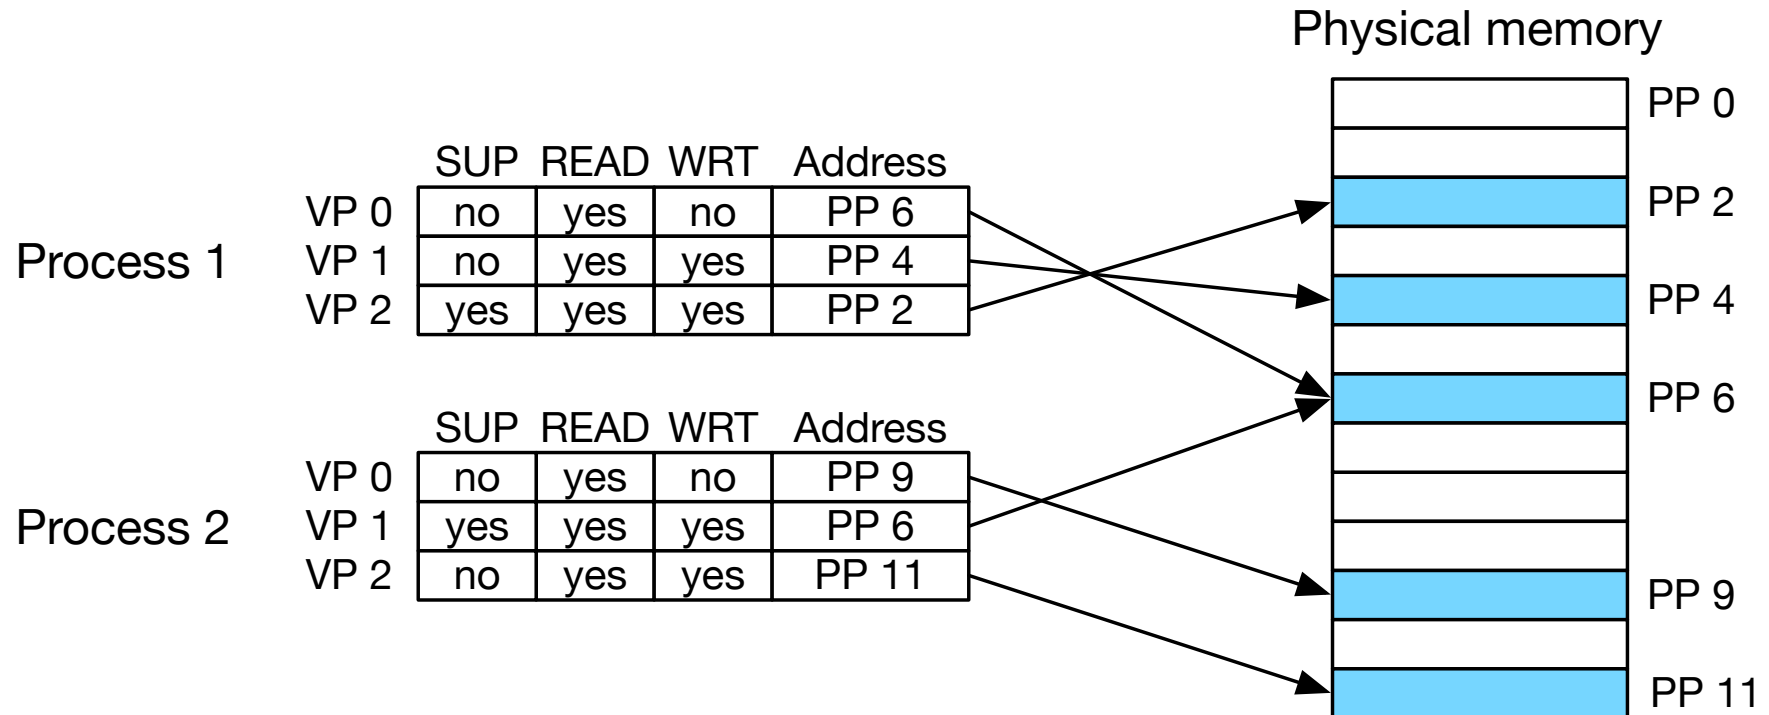
- Process may need more memory (e.g., malloc call)

⇒ New entry in page table

- Mapped to arbitrary pages in physical memory
- Do not have to be contiguous

Memory Protection

31



- Page may be kernel only: SUP=yes
- Page may be read-only (e.g., code)



address translation

Address Space

- Virtual memory size: $N = 2^n$ bytes
- Physical memory size: $M = 2^m$ bytes
- Page (block of memory): $P = 2^p$ bytes
- A virtual address can be encoded in n bits

Address Translation

- Task: mapping virtual address to physical address
 - virtual address (VA): used by machine code instructions
 - physical address (PA): location in RAM

Address Translation

- Task: mapping virtual address to physical address
 - virtual address (VA): used by machine code instructions
 - physical address (PA): location in RAM

- Formally

$$\text{MAP: } VA \rightarrow PA \cup \emptyset$$

where:

$$\begin{aligned} \text{MAP}(A) &= PA \text{ if in RAM} \\ &= \emptyset \text{ otherwise} \end{aligned}$$

Address Translation

- Task: mapping virtual address to physical address
 - virtual address (VA): used by machine code instructions
 - physical address (PA): location in RAM

- Formally

$$\text{MAP: } VA \rightarrow PA \cup \emptyset$$

where:

$$\begin{aligned} \text{MAP}(A) &= PA \text{ if in RAM} \\ &= \emptyset \text{ otherwise} \end{aligned}$$

- Note: this happens very frequently in machine code

Address Translation

- Task: mapping virtual address to physical address
 - virtual address (VA): used by machine code instructions
 - physical address (PA): location in RAM

- Formally

$$\text{MAP: } VA \rightarrow PA \cup \emptyset$$

where:


$$\begin{aligned} \text{MAP}(A) &= PA \text{ if in RAM} \\ &= \emptyset \text{ otherwise} \end{aligned}$$

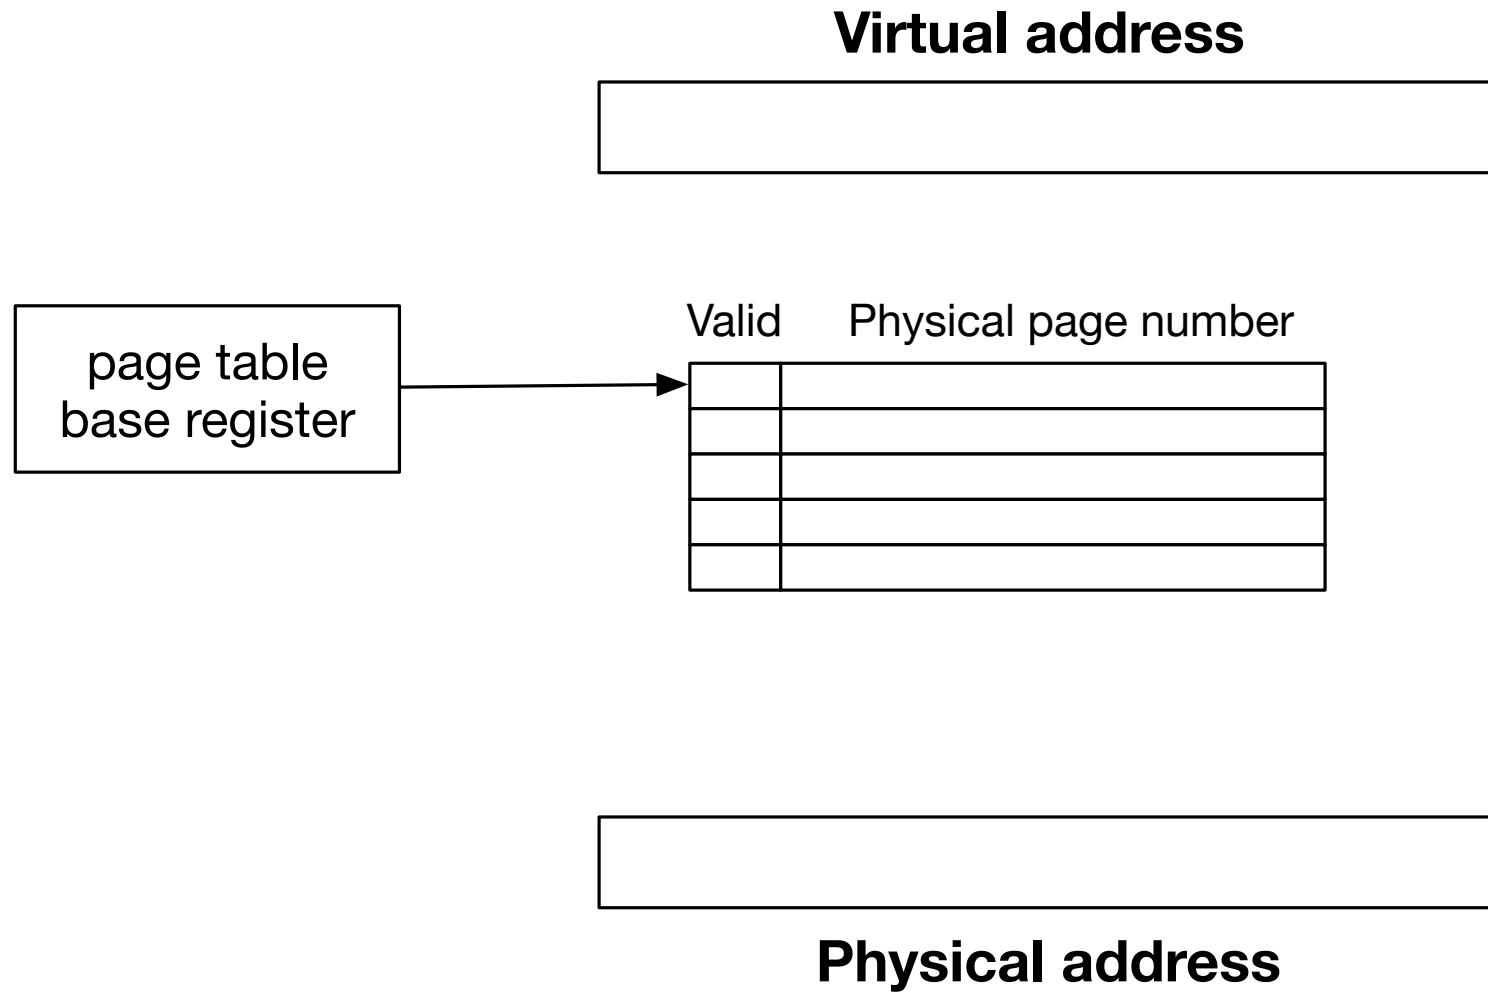
- Note: this happens very frequently in machine code
- We will do this in hardware: Memory Management Unit (MMU)

Virtual address




Physical address





Virtual address

virtual page number	page offset
---------------------	-------------



physical page number	page offset
----------------------	-------------

Physical address

