

---

# Virtual Memory

Philipp Koehn

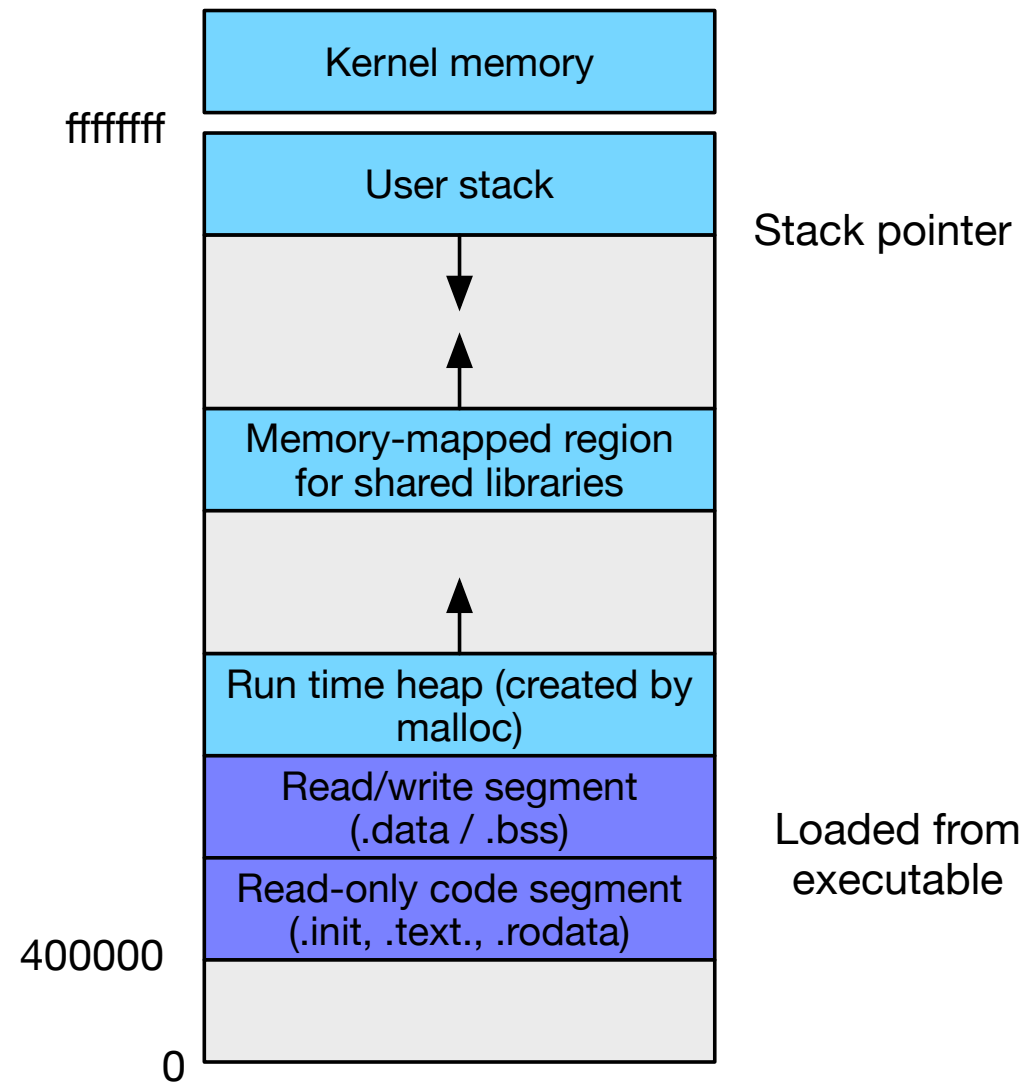
25 April 2018



# Recall: Process Address Space



1



# Virtual Memory



2

- Abstraction of physical memory
- Purpose
  - appearance of more available memory than physically exists (DRAM)
  - handles disk caching / loading
  - insulates memory of each process

# Virtual Memory



2

- Abstraction of physical memory
- Purpose
  - appearance of more available memory than physically exists (DRAM)
  - handles disk caching / loading
  - insulates memory of each process
- Page table: maps from virtual address to physical addresses

# Virtual Memory



- Abstraction of physical memory
- Purpose
  - appearance of more available memory than physically exists (DRAM)
  - handles disk caching / loading
  - insulates memory of each process
- Page table: maps from virtual address to physical addresses
- Memory management unit (MMU):  
hardware implementation of address translation


# Warning



- This is going to get very complex
- Closely tied with multi-tasking (multiple processes)
- Partly managed by hardware,  
partly managed by software

# virtual addressing

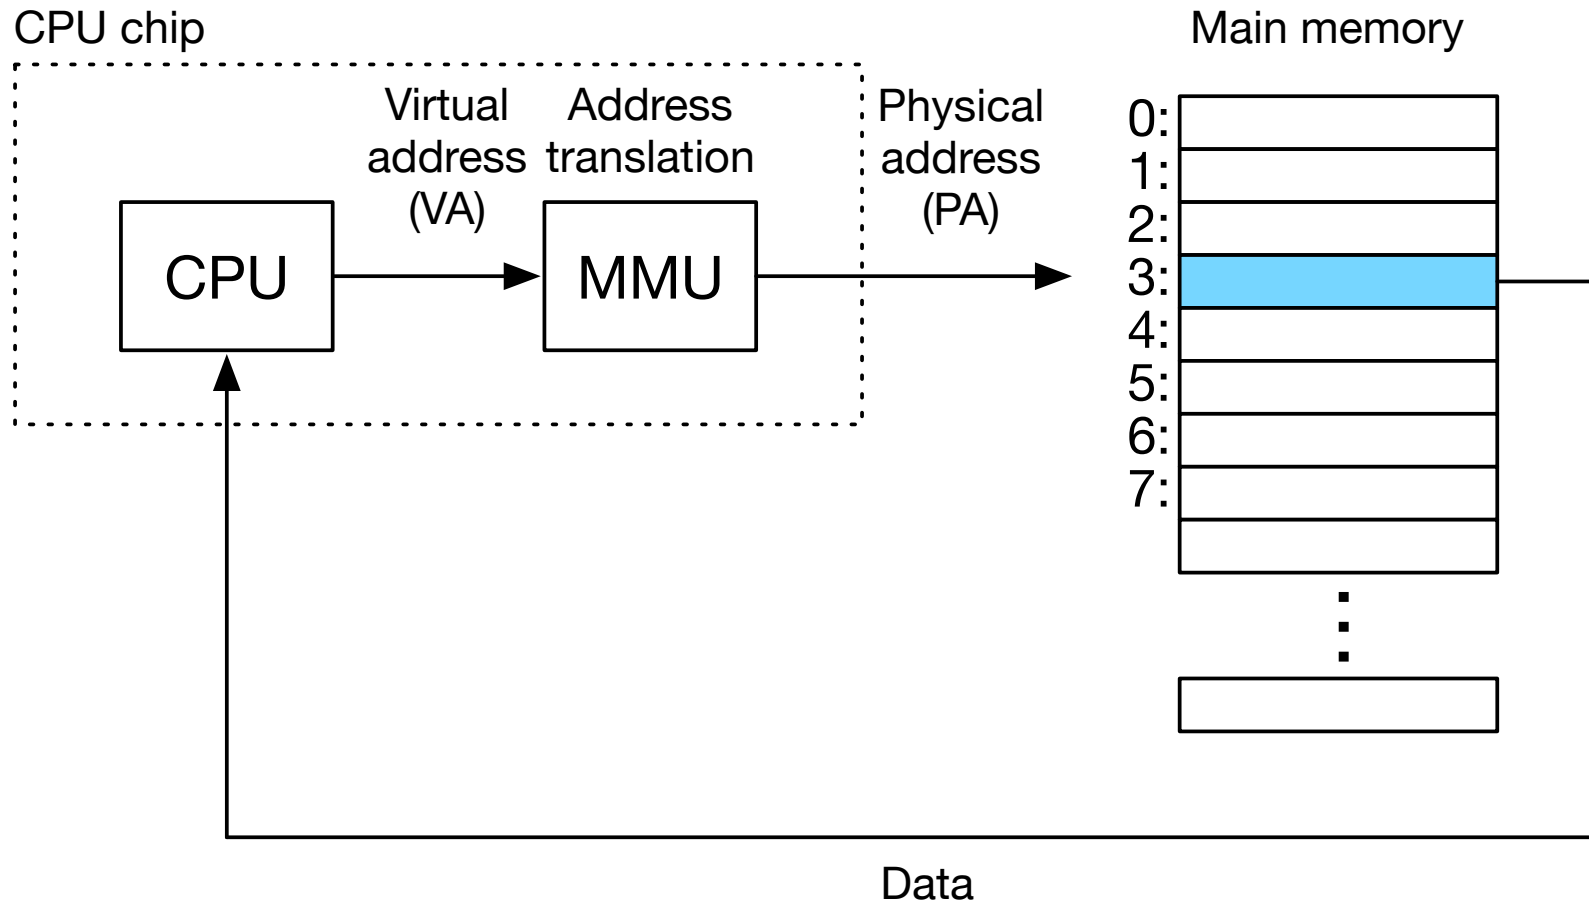
# Physical Addressing



- So far, assumed CPU addresses physical memory



# Virtual Addressing



- Memory management unit (MMU): maps virtual to physical addresses

# Address Space



- Virtual memory size:  $N = 2^n$  bytes, e.g., 256TB
- Physical memory size:  $M = 2^m$  bytes, e.g., 16GB
- Page (block of memory):  $P = 2^p$  bytes, e.g., 4KB
- A virtual address can be encoded in  $n$  bits

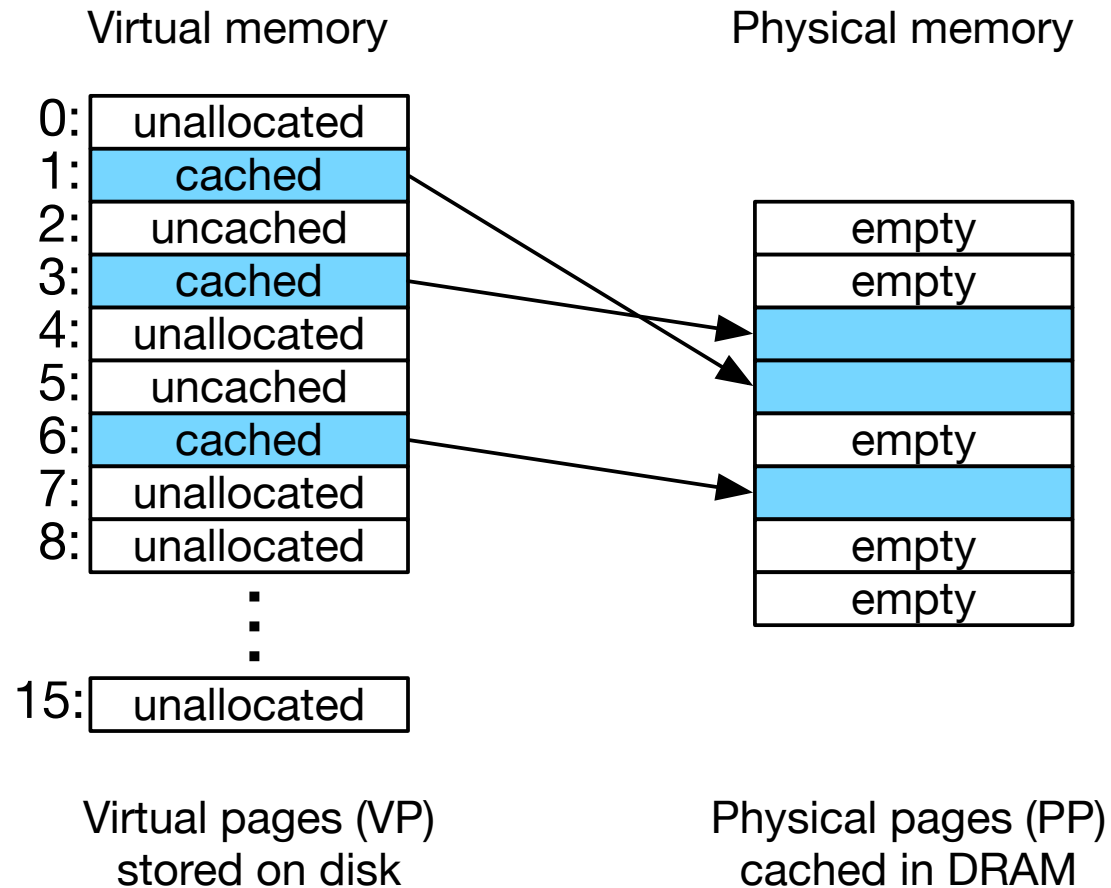
# caching

# Caching... Again?



- Yes, we already discussed caching  
but: for on-chip cache of DRAM memory
- Now
  - caching between RAM and disk
  - driven by a large virtual memory address space
  - to avoid unnecessary and duplicate loading
- Jargon
  - previously "block", now "page"
  - now: "swapping" or "paging"

# Mapping



# State of Virtual Memory Page

11



- Cached
  - allocated page
  - stored in physical memory

# State of Virtual Memory Page

- Cached
  - allocated page
  - stored in physical memory
- Uncached
  - allocated page
  - not in physical memory

# State of Virtual Memory Page

- Cached
  - allocated page
  - stored in physical memory
- Uncached
  - allocated page
  - not in physical memory
- Unallocated
  - not used by virtual memory system so far



# Page Table

12



- Array of page table entries (PTE)

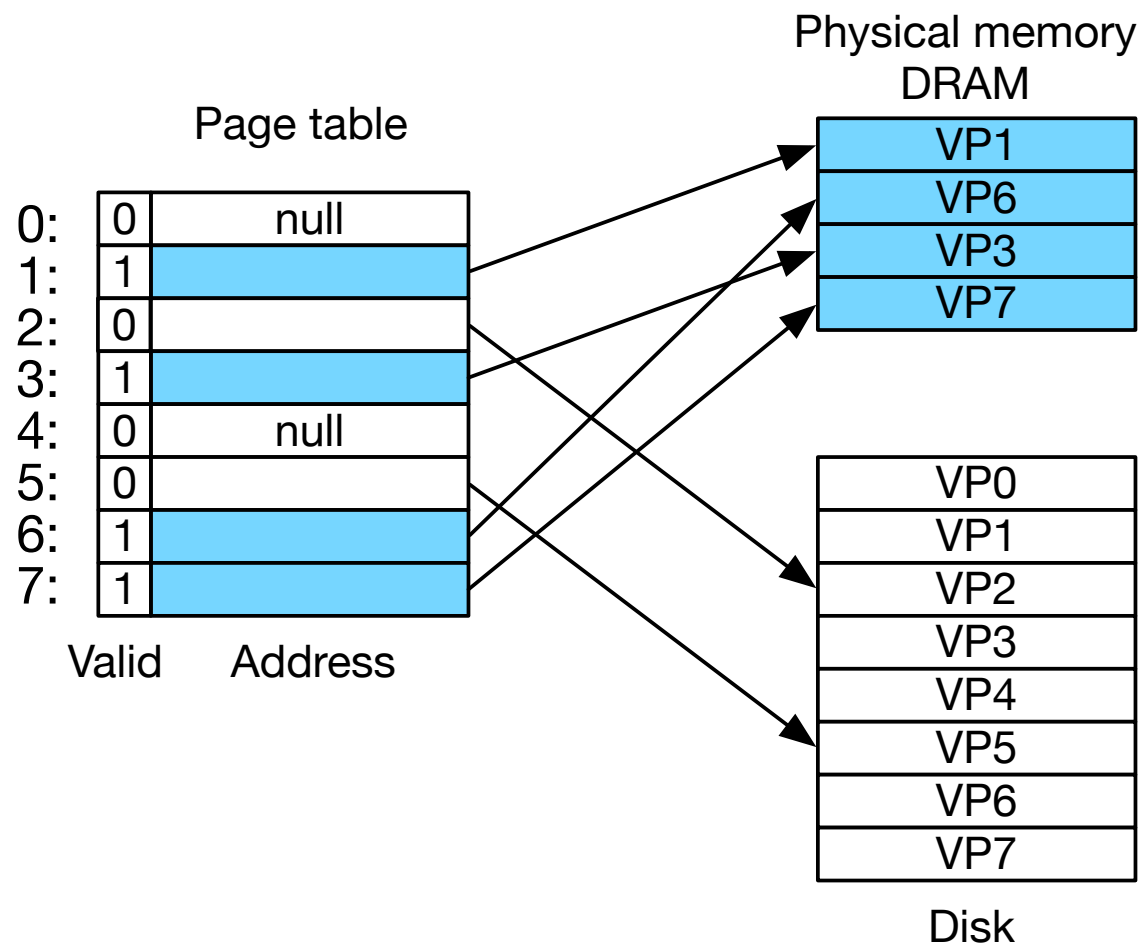
# Page Table

- Array of page table entries (PTE)
- Valid bit
  - set if PTE currently maps to physical address (cached)
  - not set otherwise (uncached or unallocated)

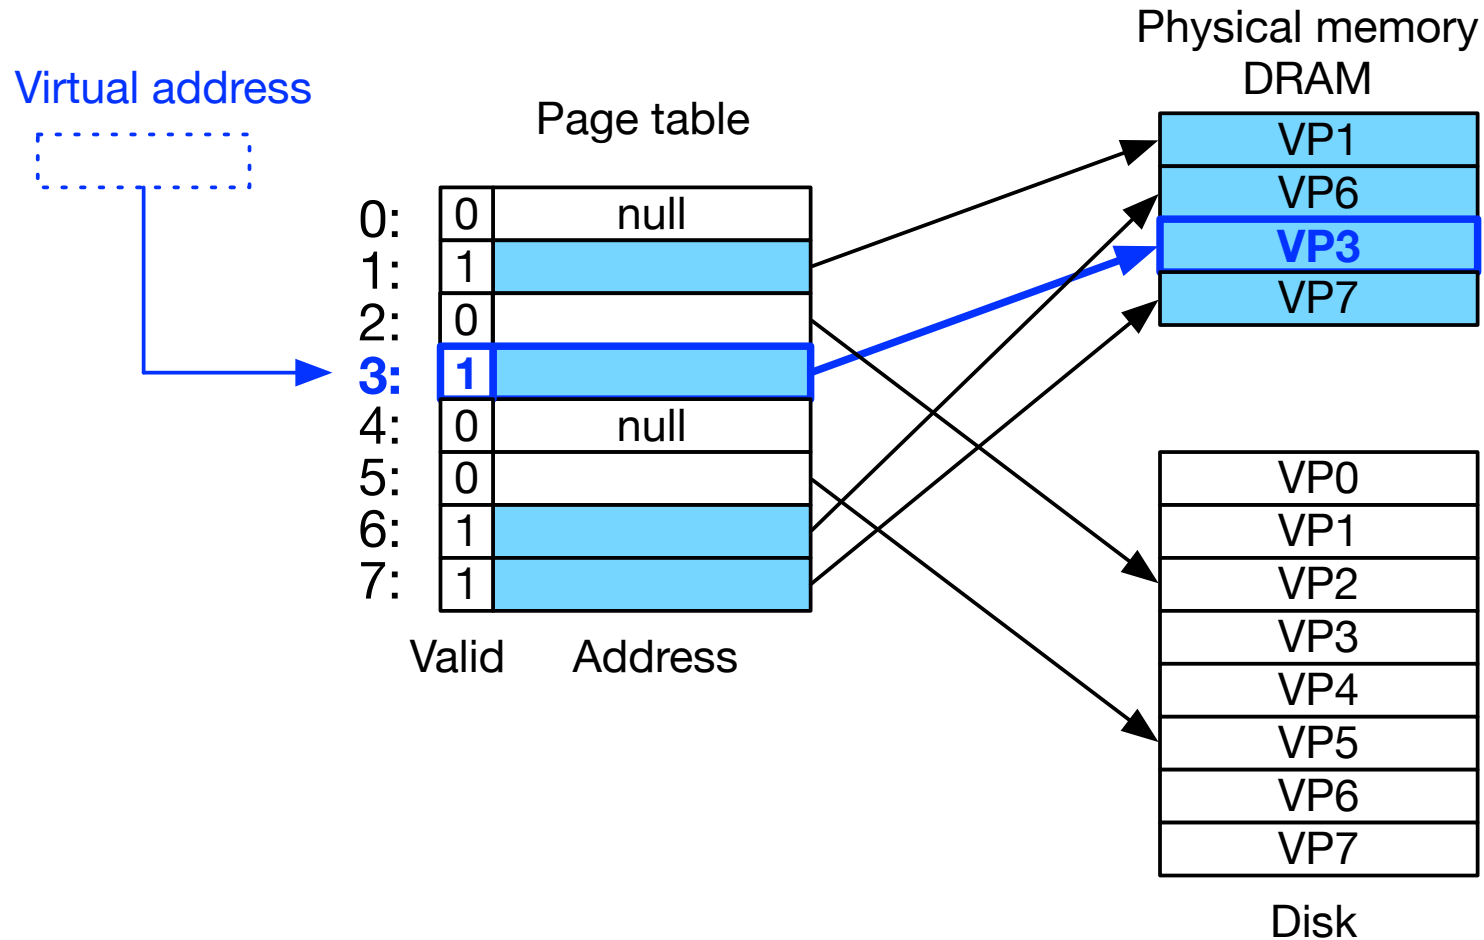
# Page Table

- Array of page table entries (PTE)
- Valid bit
  - set if PTE currently maps to physical address (cached)
  - not set otherwise (uncached or unallocated)
- Mapped address
  - if cached: physical address in DRAM
  - if not cached: physical address on disk

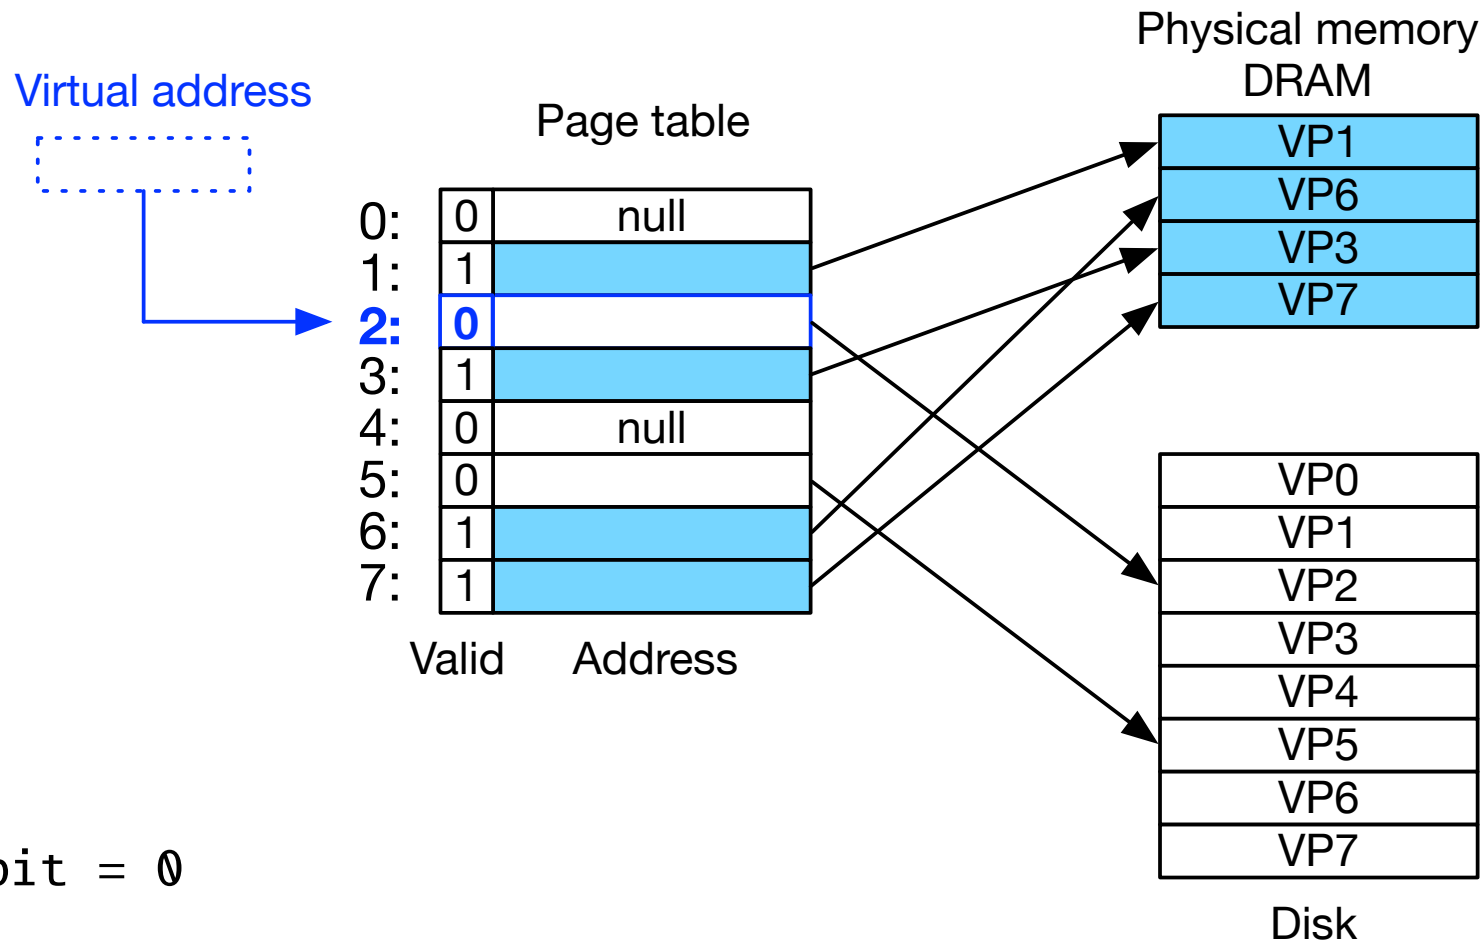
# Page Table



# Page Hit



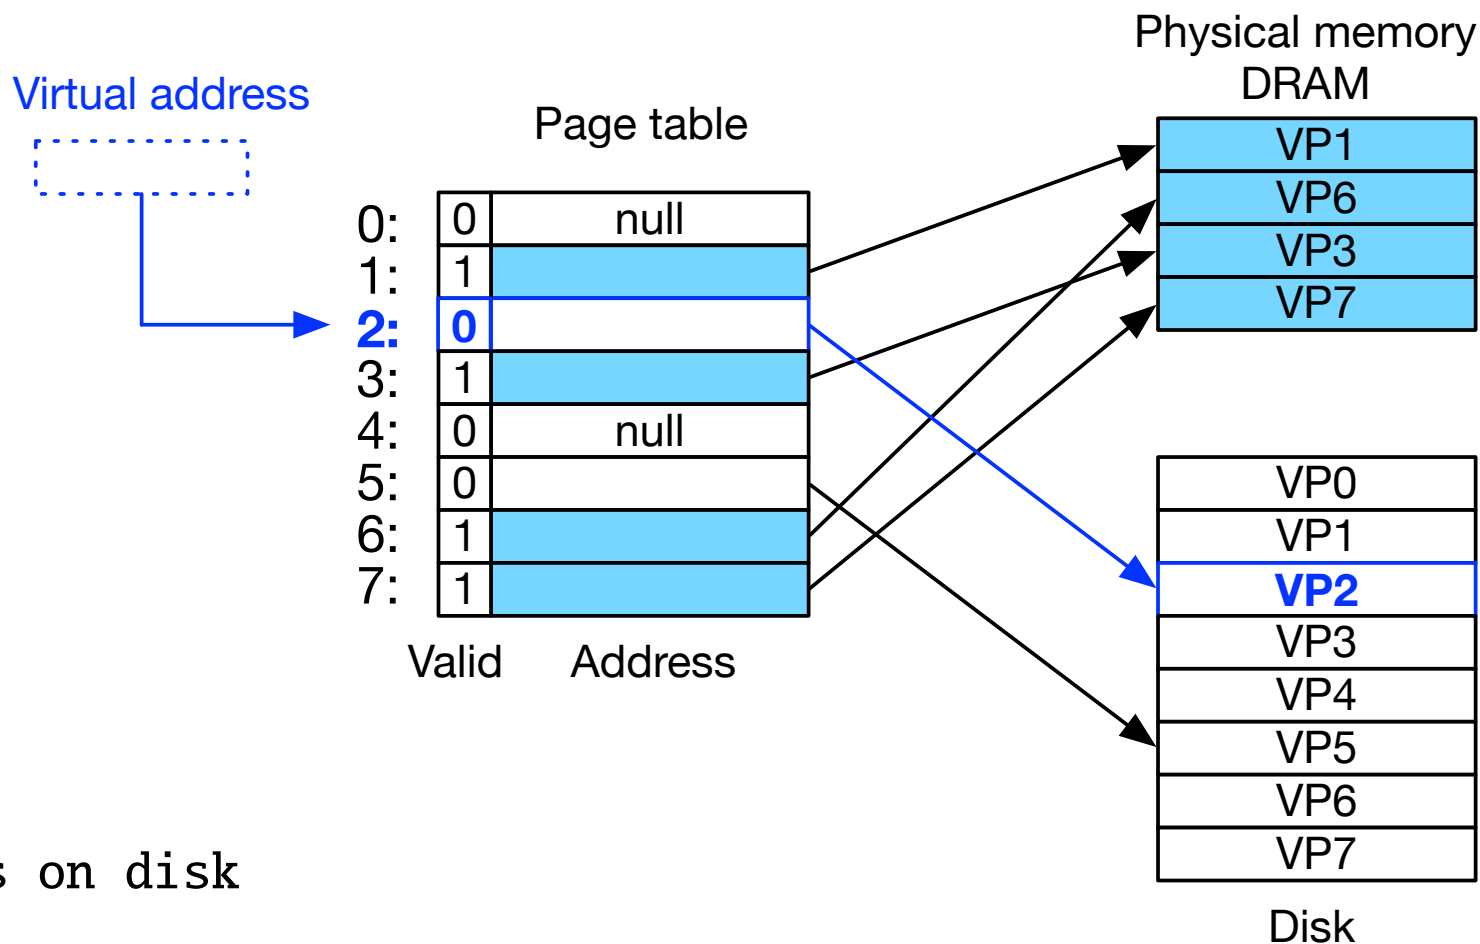
# Page Fault



- Valid bit = 0
- Page not in RAM

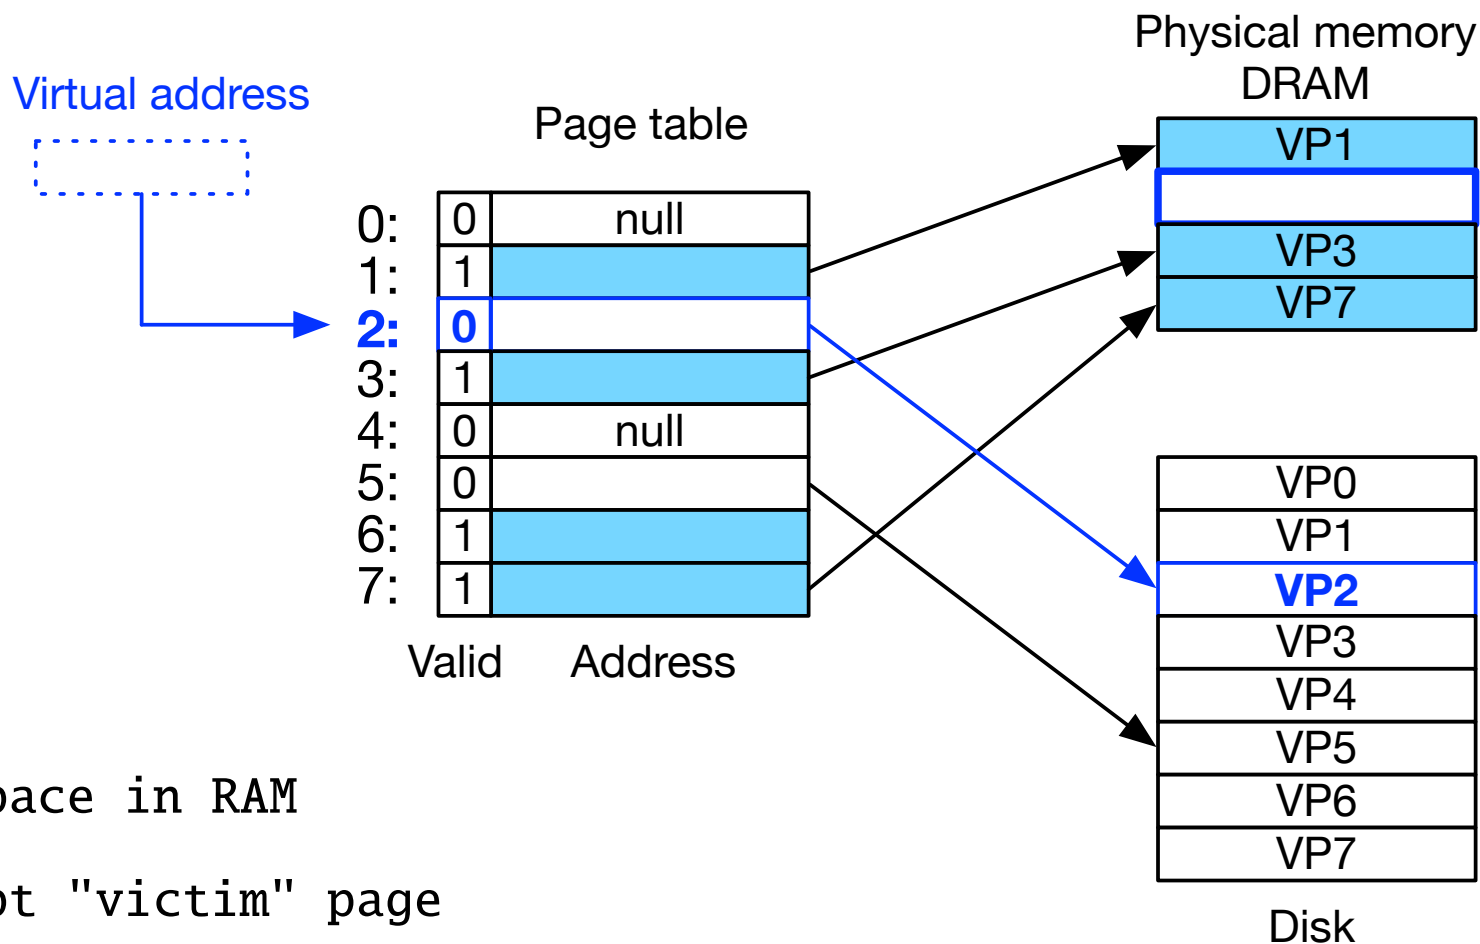
# Page Fault

16



- Page is on disk

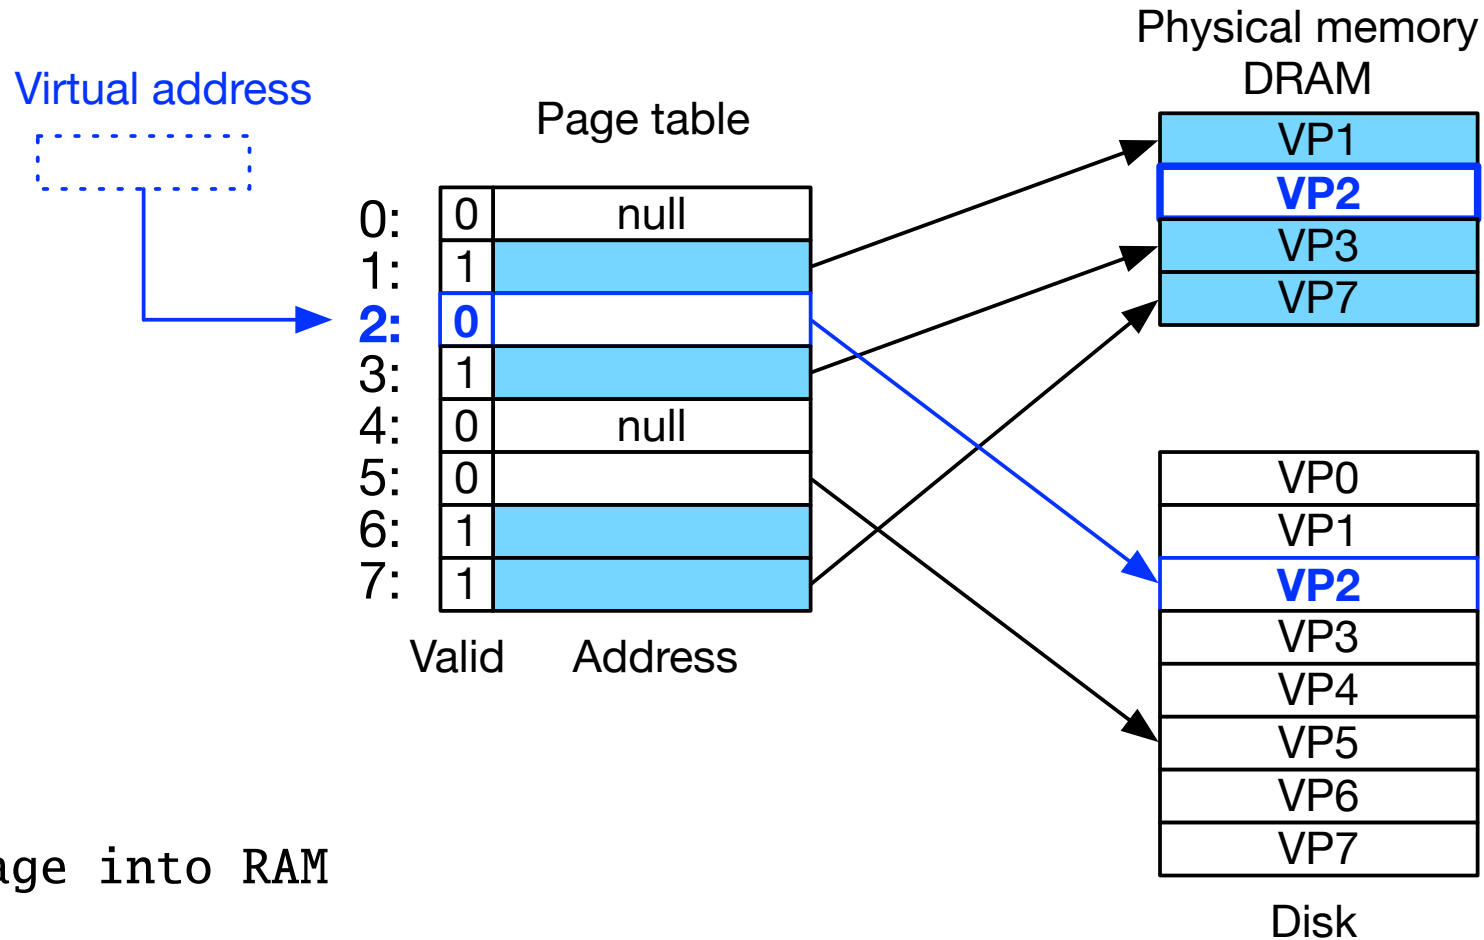
# Page Fault



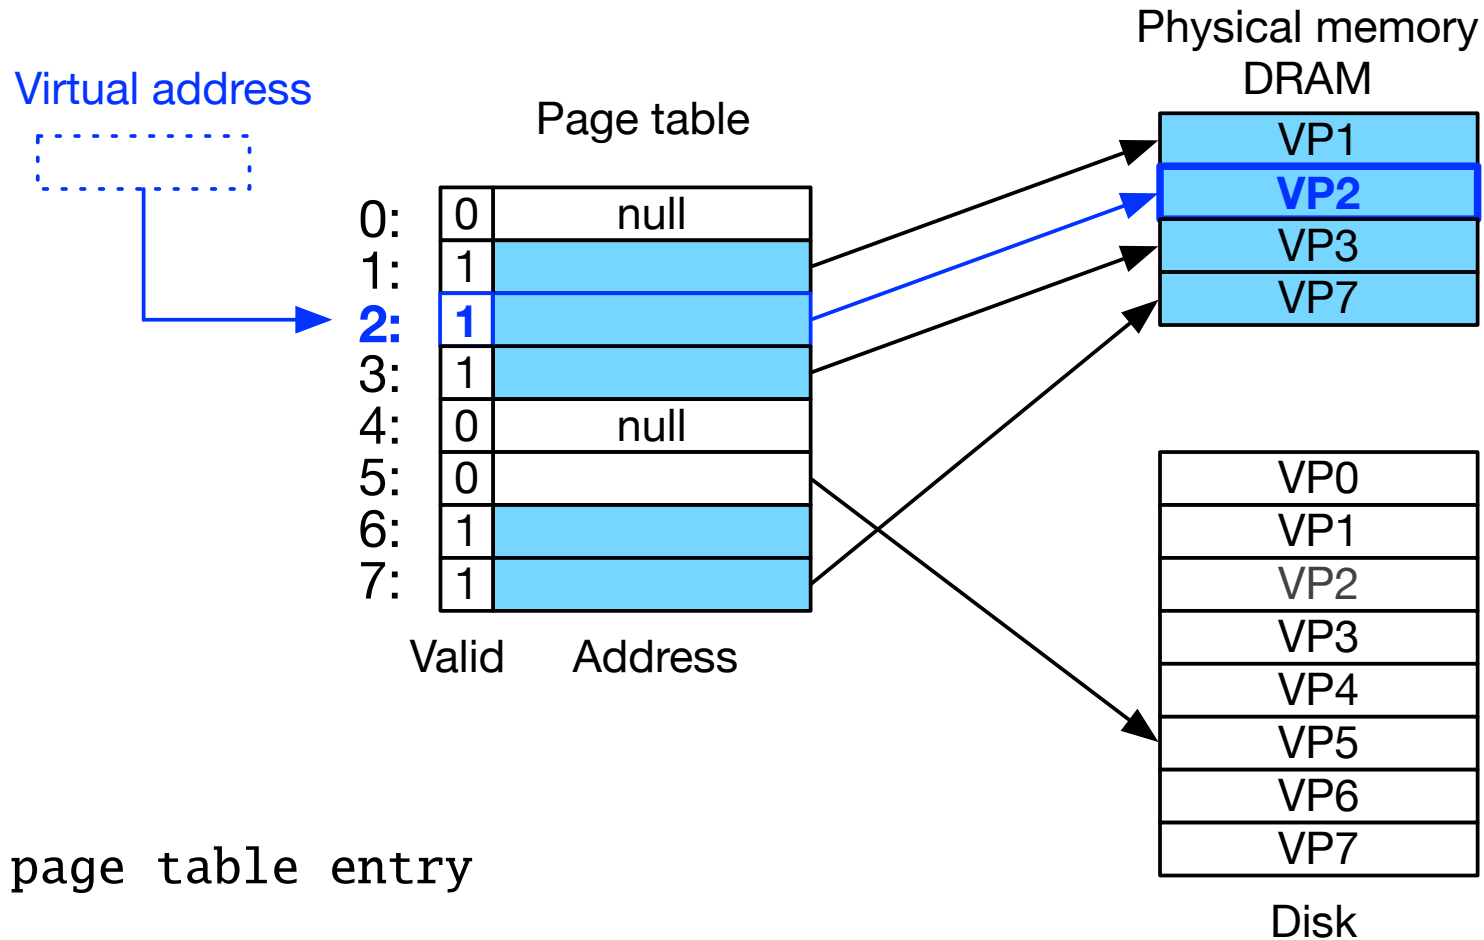
- Make space in RAM
- Pre-empt "victim" page
- Typically out-dated cached page



# Page Fault



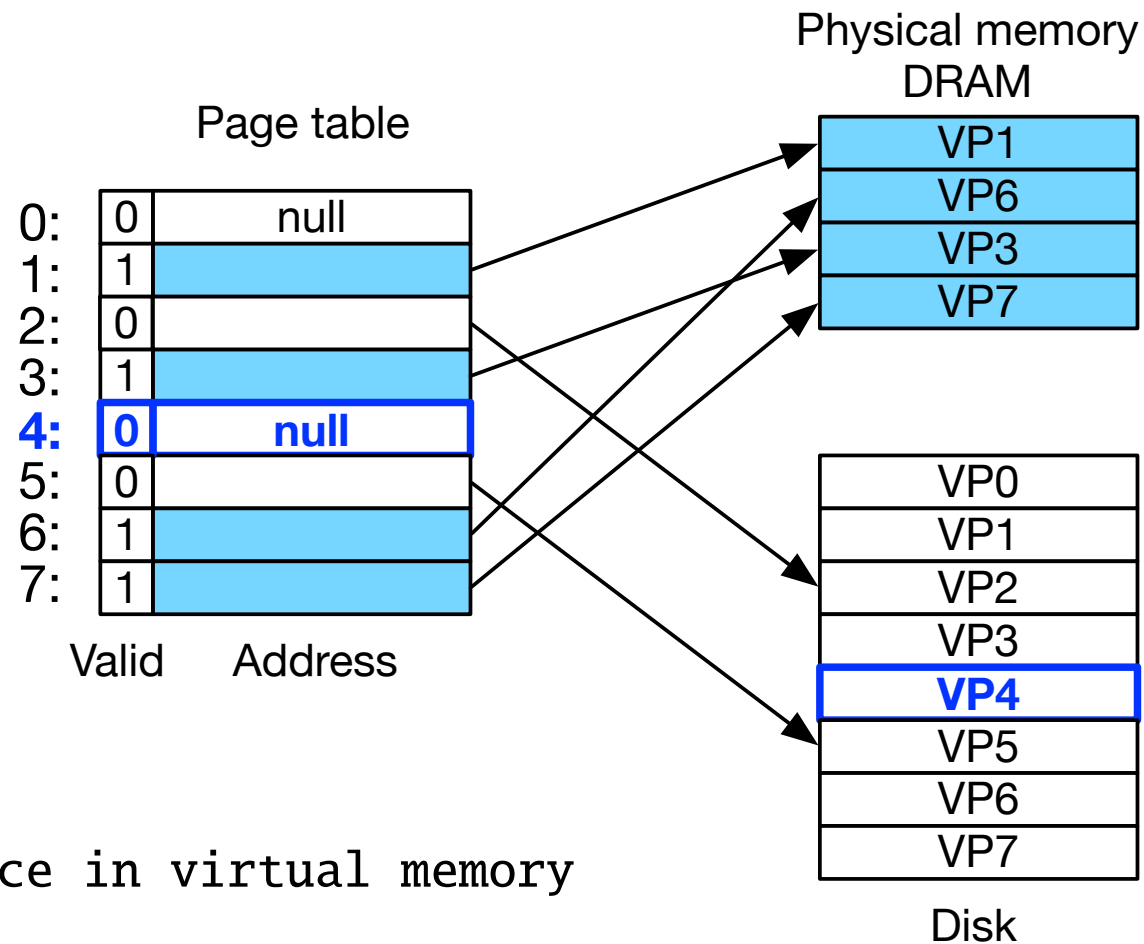
# Page Fault



# Allocating Pages

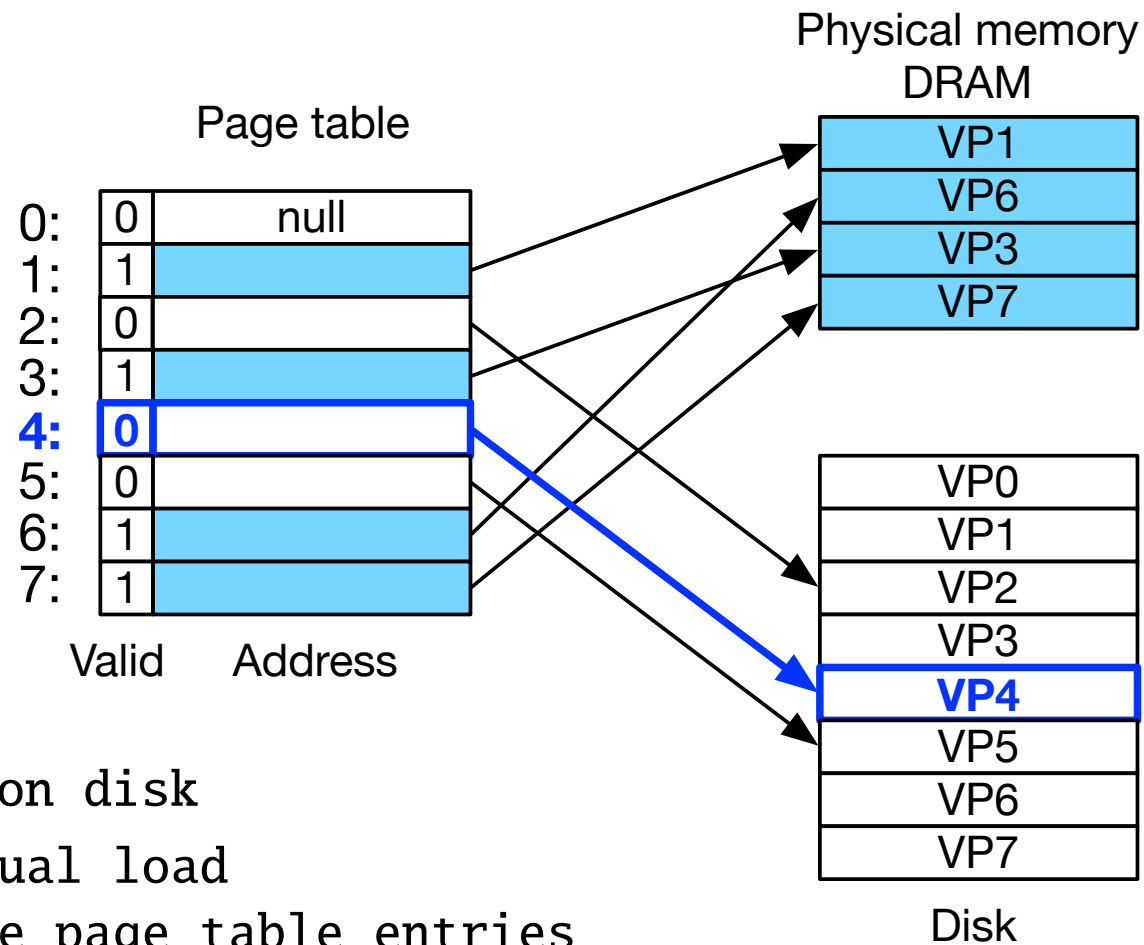
- What happens when we load a program?
- We need to load its executable into memory
- Similar: create data objects when program is running ("allocating" memory)

# Allocating Page



- Identify space in virtual memory

# Allocating Page



- Map to data on disk
  - do not actual load
  - just create page table entries
  - let virtual memory system handle loading

⇒ On-demand loading

# Process Memory

- Nothing loaded at startup

# Process Memory

- Nothing loaded at startup
- Working set (or resident set)
  - pages of a process that are currently in DRAM
  - loaded by virtual memory system on demand

# Process Memory

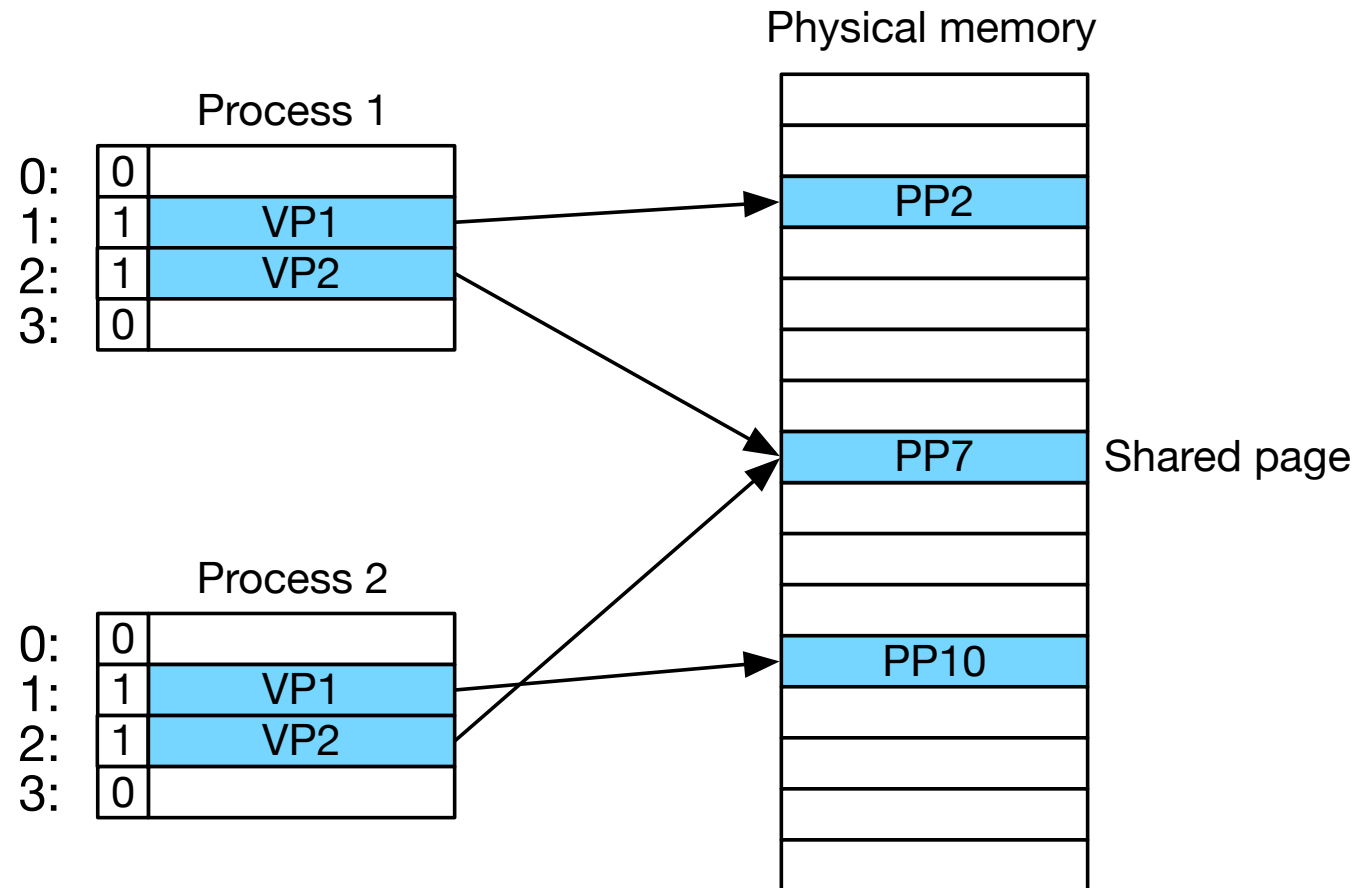
- Nothing loaded at startup
- Working set (or resident set)
  - pages of a process that are currently in DRAM
  - loaded by virtual memory system on demand
- Thrashing
  - memory actively required by all processes larger than physically available
  - frequent swapping of memory to/from disk
  - very bad: slows down machine dramatically



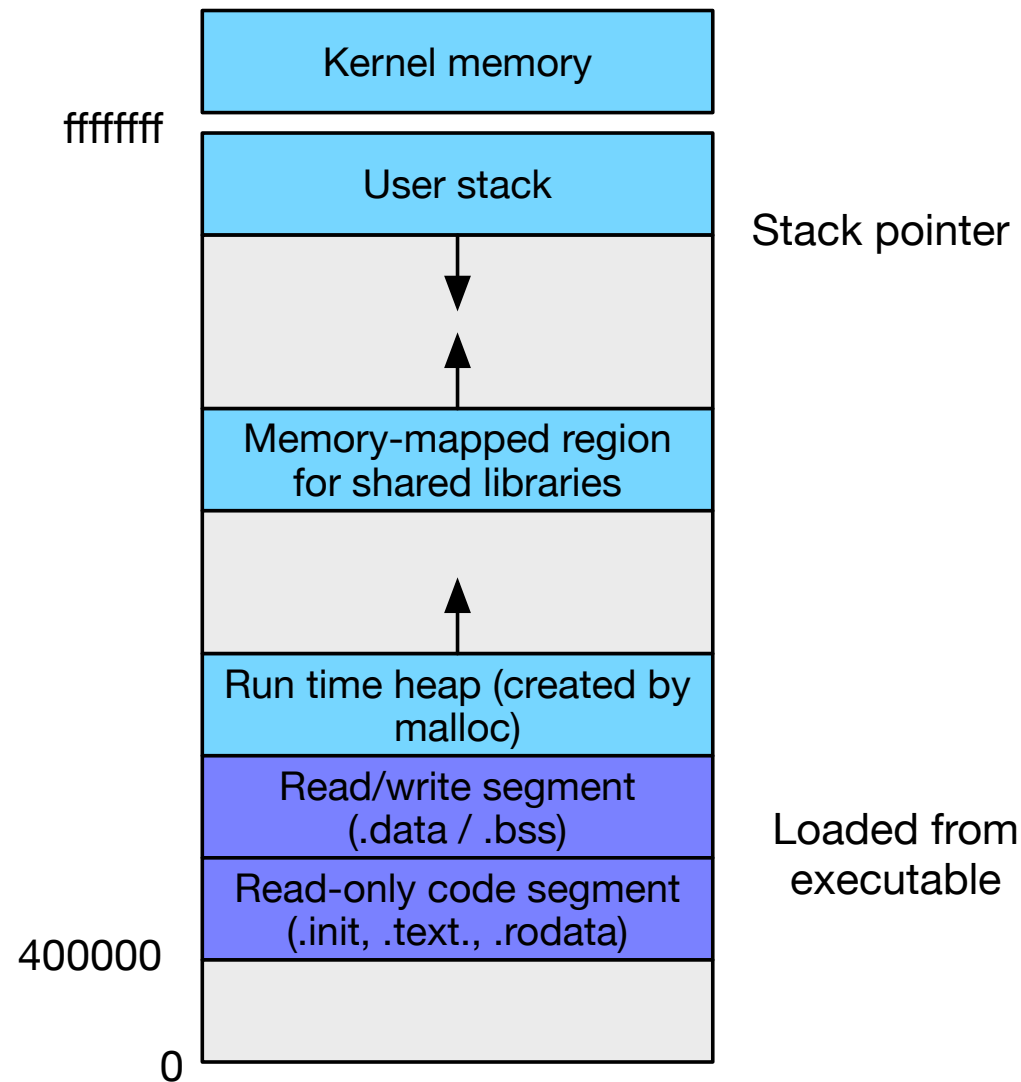


# memory management


# One Page Table per Process



# Process Address Space



# Simplified Linking



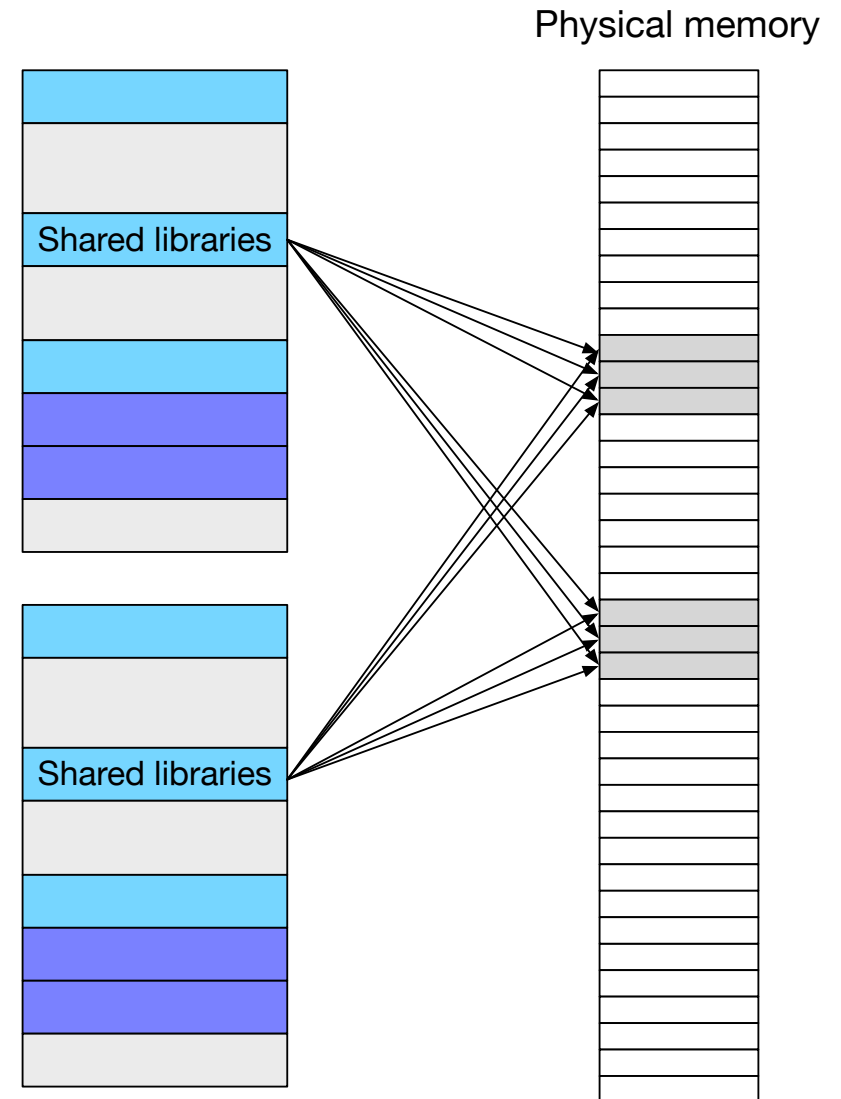
- Each process has its code in address `0x400000`
- Easy linking: Linker can establish fixed addresses

# Simplified Loading

- When loading process into memory...
- Enter .data and .text section into page table
- Mark them as invalid (= not actually in RAM)
- Called memory mapping (more on that later)

# Simplified Sharing

- Shared libraries used by several processes  
e.g., `stdio` providing `printf`, `scanf`, `open`, `close`, ...
- Not copied multiple times into RAM



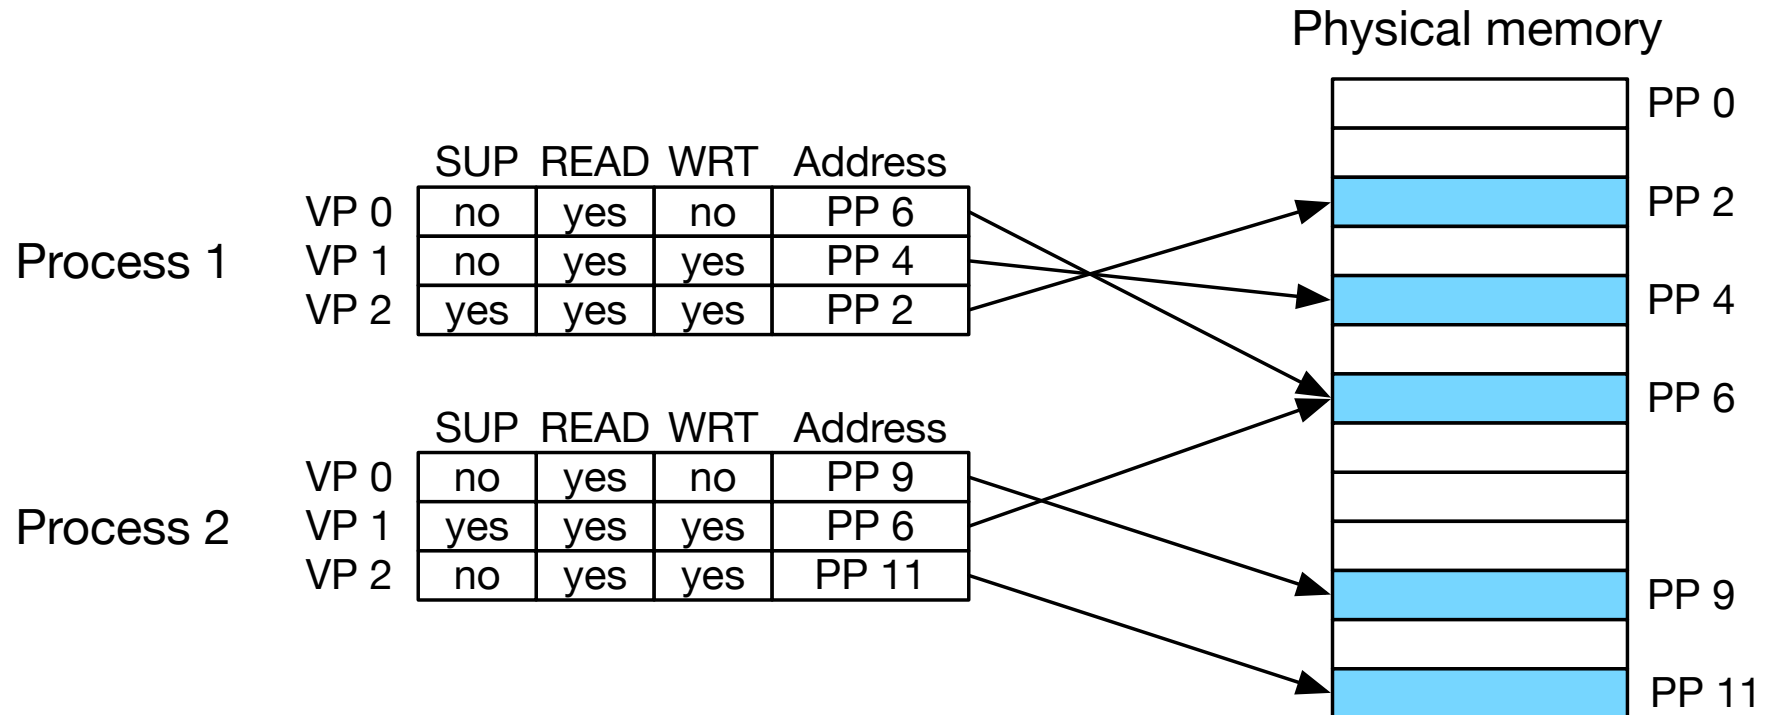
# Simplified Memory Allocation

- Process may need more memory (e.g., malloc call)

⇒ New entry in page table

- Mapped to arbitrary pages in physical memory
- Do not have to be contiguous

# Memory Protection



- Page may be kernel only: SUP=yes
- Page may be read-only (e.g., code)





# address translation

# Address Space

- Virtual memory size:  $N = 2^n$  bytes
- Physical memory size:  $M = 2^m$  bytes
- Page (block of memory):  $P = 2^p$  bytes
- A virtual address can be encoded in  $n$  bits

# Address Translation

- Task: mapping virtual address to physical address
  - virtual address (VA): used by machine code instructions
  - physical address (PA): location in RAM

# Address Translation

- Task: mapping virtual address to physical address
  - virtual address (VA): used by machine code instructions
  - physical address (PA): location in RAM

- Formally

$$\text{MAP: } VA \rightarrow PA \cup \emptyset$$

where:

$$\begin{aligned} \text{MAP}(A) &= PA \text{ if in RAM} \\ &= \emptyset \text{ otherwise} \end{aligned}$$

# Address Translation

- Task: mapping virtual address to physical address
  - virtual address (VA): used by machine code instructions
  - physical address (PA): location in RAM

- Formally

$$\text{MAP: } VA \rightarrow PA \cup \emptyset$$

where:

$$\begin{aligned} \text{MAP}(A) &= PA \text{ if in RAM} \\ &= \emptyset \text{ otherwise} \end{aligned}$$

- Note: this happens very frequently in machine code

# Address Translation

- Task: mapping virtual address to physical address
  - virtual address (VA): used by machine code instructions
  - physical address (PA): location in RAM

- Formally

$$\text{MAP: } VA \rightarrow PA \cup \emptyset$$

where:


$$\begin{aligned} \text{MAP}(A) &= PA \text{ if in RAM} \\ &= \emptyset \text{ otherwise} \end{aligned}$$

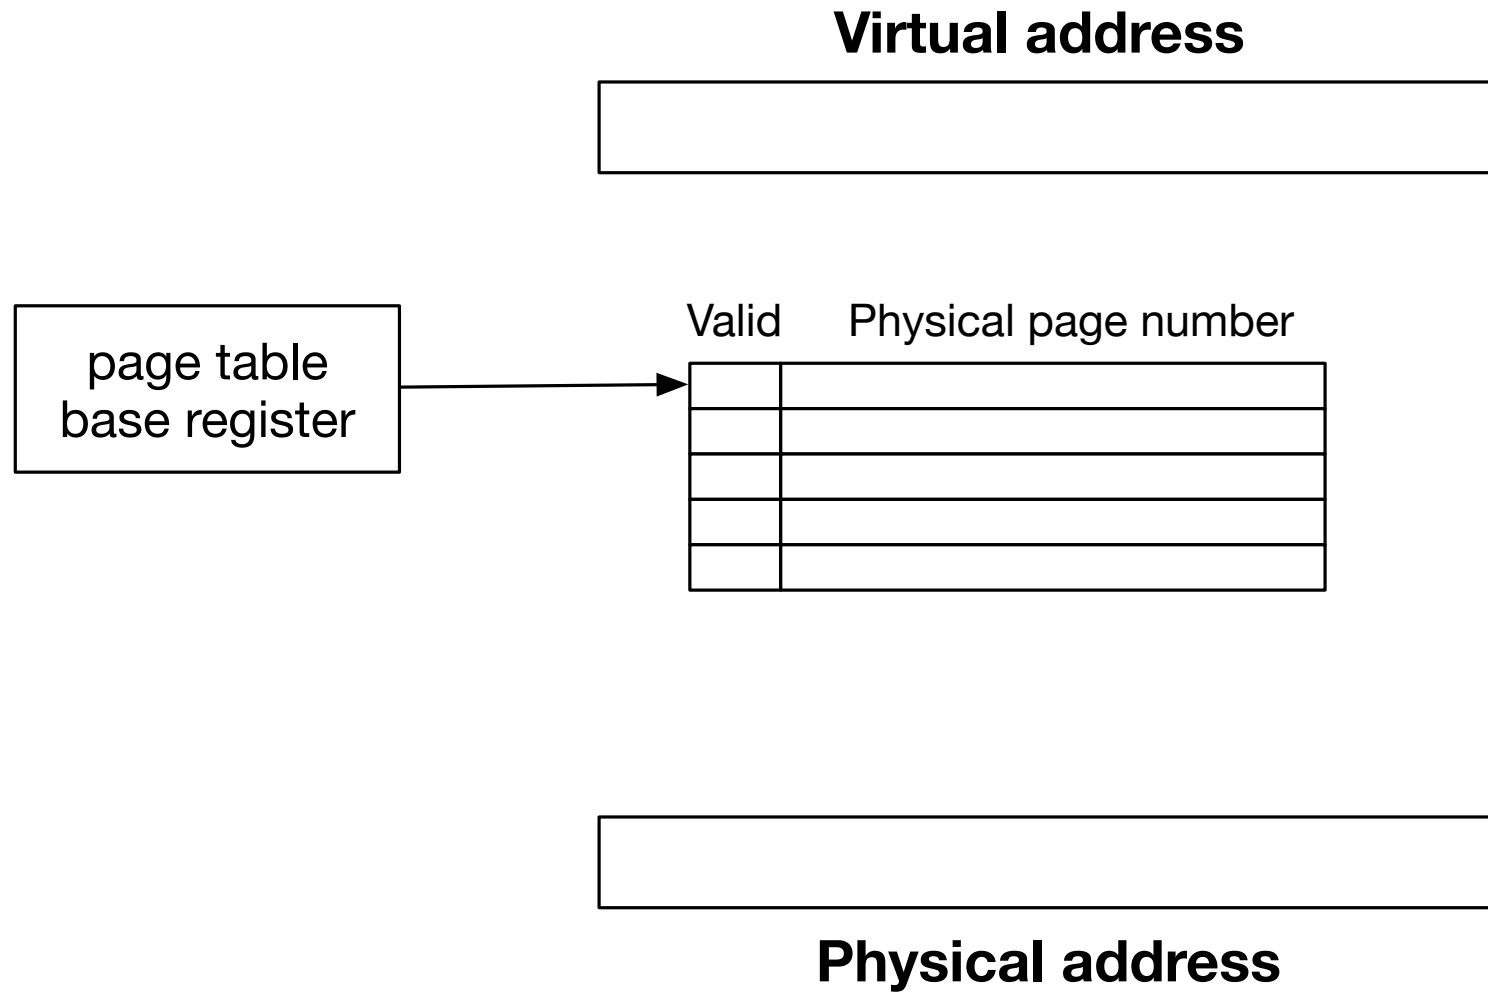
- Note: this happens very frequently in machine code
- We will do this in hardware: Memory Management Unit (MMU)

**Virtual address**



**Physical address**








## Virtual address

virtual page number	page offset
---------------------	-------------



physical page number	page offset
----------------------	-------------

## Physical address

